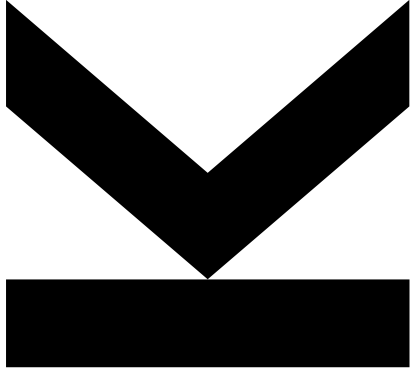


Trees (Height Balanced)



Algorithms and Data Structures 2, 340300
Lecture – 2023W
Univ.-Prof. Dr. Alois Ferscha, teaching@pervasive.jku.at

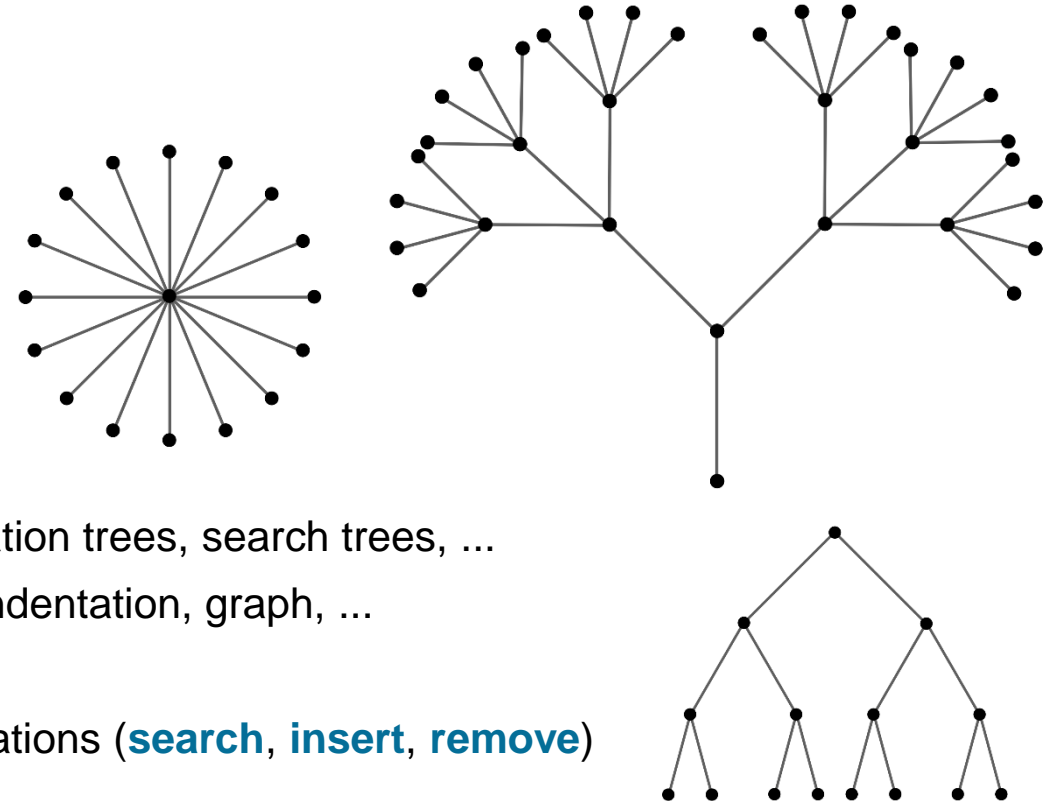
Definition & Terminology

A **tree** is

- an **acyclic, simple, coherent** graph
- i.e. does **not** contain **loops and cycles**:
between each pair of nodes there is at maximum one edge

Generalization of lists:

- Element (node) may have **multiple successors**.
- Exactly 1 node has no predecessor: **root**
- Nodes without successors: **leaves**



Frequently used data structure: decision trees, syntax trees, derivation trees, search trees, ...

Frequently used representation form: quantity, bracket, recursive indentation, graph, ...

Here: Use of trees to **store keys** and realization of dictionary operations (**search**, **insert**, **remove**) in e.g. binary trees

Definition & Terminology

Tree B is **ordered**, if successors of each node are ordered (1., 2., 3. etc.; left, right).

In an ordered tree, the **subtrees B_i** of each node form an **ordered set** (e.g.: arithmetic expression)

Order of B : maximum number of **successors** of a node

Path of length k : Follow p_0, \dots, p_k of nodes, such that p_i is successor of p_{i-1}

- **Height of a tree**:
maximum distance of a leaf to the root.
- **Depth of a node**:
distance to the root (the **number of edges** on a path from this node to the root)
The nodes at level i are all nodes with depth i .

A tree of order n is called **complete** if all leaves have the **same depth** and the maximum number of nodes is present at each level.



Definition & Terminology

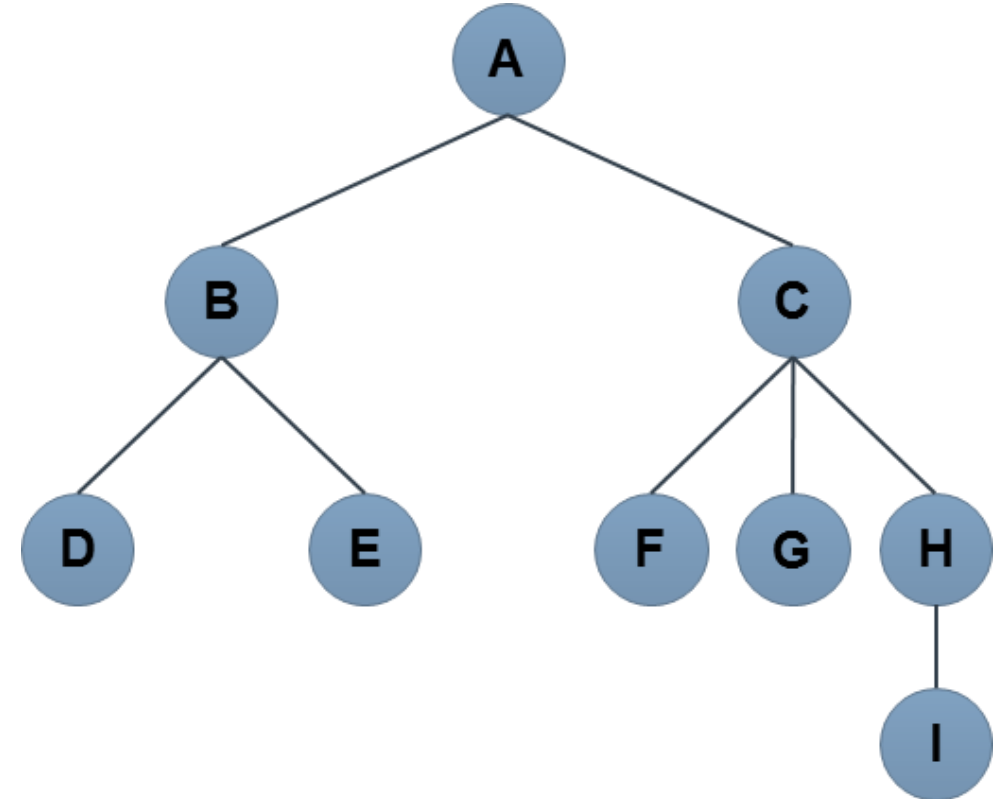
A	root
B	is parent of D and E
C	is sibling of B
D and E	are children of B
D, E, F, G, I	are external nodes or leaves
A, B, C, H	are internal nodes

The **depth** of E is 2.

The **height** of the tree is 3.

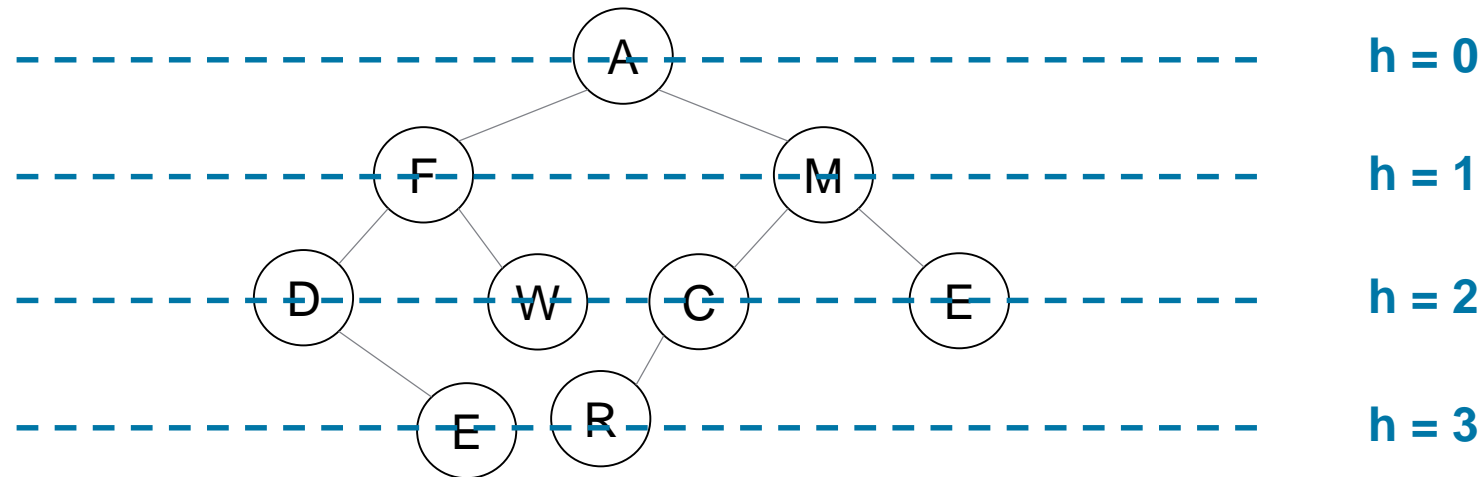
The **order** of B is 2, the **order** of C is 3.

We always have: number of edges = number of nodes - 1



Binary Trees

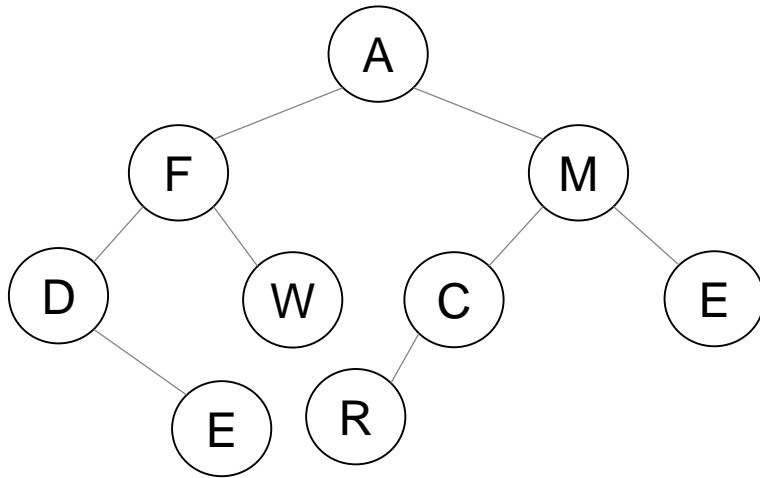
If T is a **binary tree** (order of all nodes ≤ 2) with **n nodes** and a **height h** , then T has the following properties



- Number of **external nodes** is at least $h+1$ and at maximum 2^h

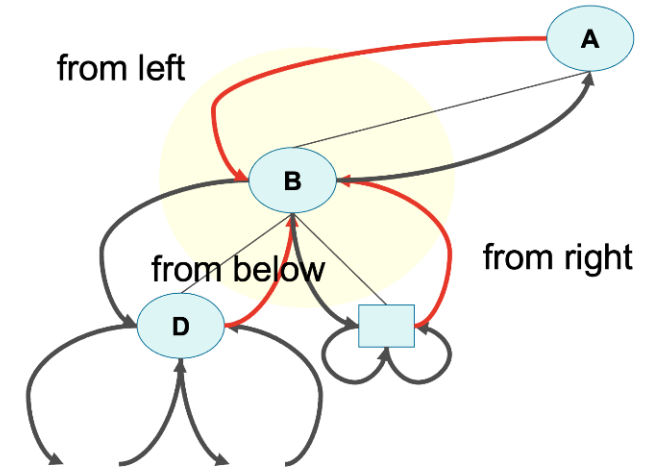
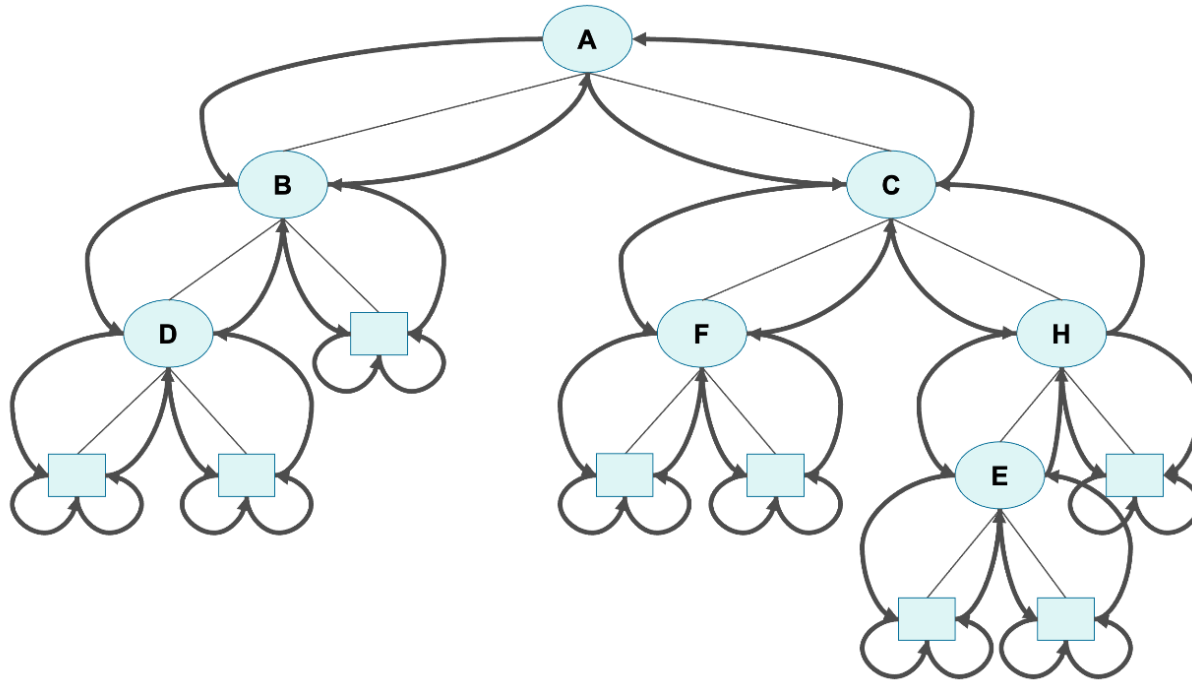
Binary Trees

If T is a **binary tree** (order of all nodes ≤ 2) with n nodes and a height h , then T has the following properties

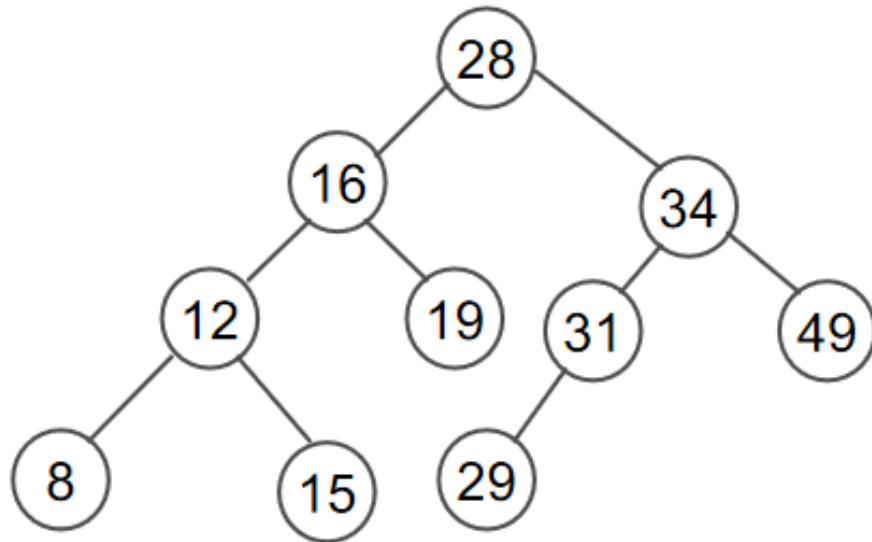


- Number of **external nodes** is at least $h+1$ and at maximum 2^h
- Number of **internal nodes** is at least h and at maximum $2^h - 1$
- Number of **nodes** is at least $2h+1$ and at maximum $2^{h+1} - 1$
- The **height** of T is at least $\log(n+1) - 1$ and at maximum $(n-1)/2$

Traversing Binary Trees



Traversing Binary Trees



- **Preorder:**
28, 16, 12, 8, 15, 19, 34, 31, 29, 49
- **Postorder:**
8, 15, 12, 19, 16, 29, 31, 49, 34, 28
- **Inorder:**
8, 12, 15, 16, 19, 28, 29, 31, 34, 49

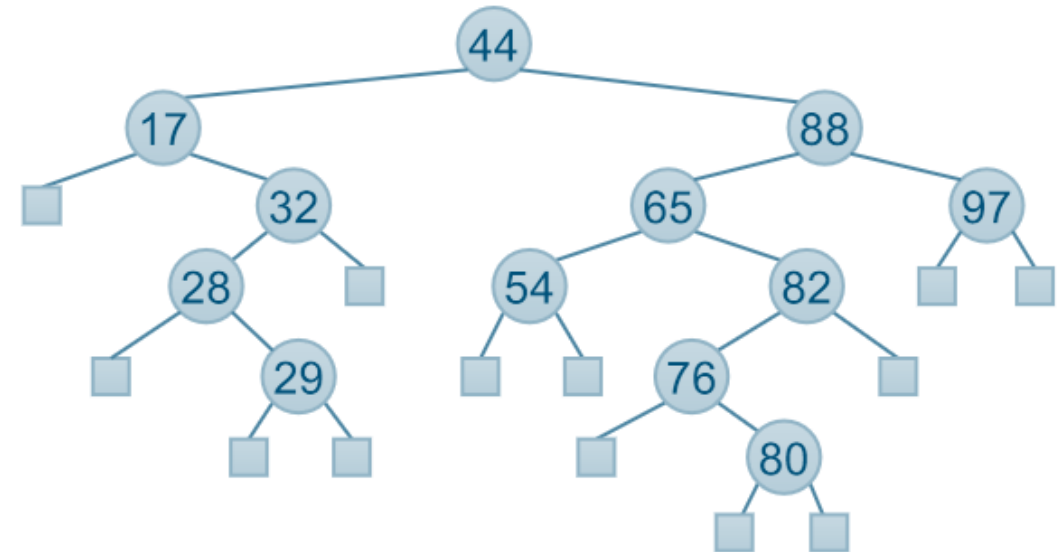
Hints: If you traverse the tree starting from the root, you have

- **Preorder** by visiting all nodes on the left side of which you pass by.
- **Postorder** by visiting all nodes on the right side of which you pass by.
- **Inorder** by visiting all nodes on the bottom side of which you pass by.

Binary Search Trees

A **binary search tree** is a **binary tree** T where:

- Each internal node stores a **key-value pair** of a dictionary
- Keys which are stored in nodes of the **left subtree** of a node v , are **less than or equal** to the key stored in v
- Keys which are stored in nodes of the **right subtree** of a node v , are **greater than** the key stored in v
- **External nodes** serve only as **placeholders** and do not store elements



Searching in Binary Search Trees

Assumption: For each node n with children c_l , c_r and key k we have (search tree condition):

- All keys stored in subtree of c_l are **smaller** than k
- All keys stored in subtree of c_r are **larger** than k
- All keys stored in subtree of c_l are **smaller than all** keys stored in subtree of c_r

Search for key x :

1. Compare key k of inspected node n with x .
2. If $n == \text{null}$, x is **not in** the binary tree.
3. If $k < x$, set $n = c_r(n)$ and jump to 1.
4. If $k > x$, set $n = c_l(n)$ and jump to 1.
5. If $k == x$, found.

Maximum number of inspected nodes: **height of tree**

Search within the nodes, e.g. by linear or binary search. As $l \leq m$, the complexity is constant.

Searching in Binary Search Trees

Binary search tree is a **decision tree**:

in each **internal node v** the key to be searched for is compared with the **key stored in v**

Pseudocode

Algorithm TreeSearch(k, v)

Input: key k to be searched, node v of a binary search tree

Output: node w in subtree of v,

for successful search internal node with key k,

for unsuccessful search

external node after all internal nodes smaller than k and before all internal nodes greater than k

if v is an external node **then**

return v

if k = key(v) **then**

return v

else if k < key(v) **then**

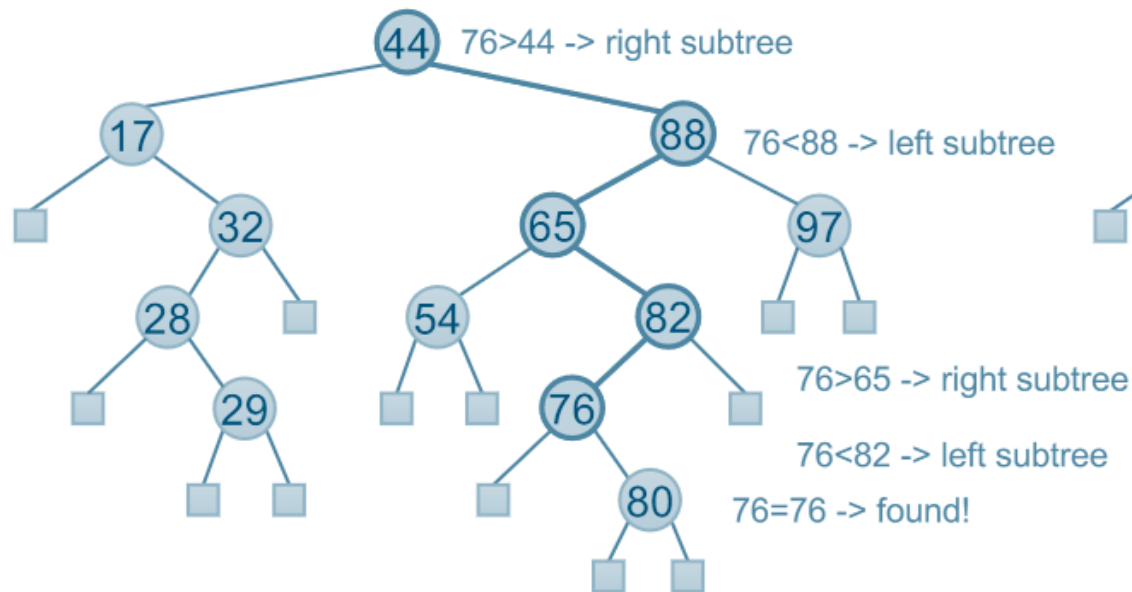
return TreeSearch(k, T.leftChild(v))

else

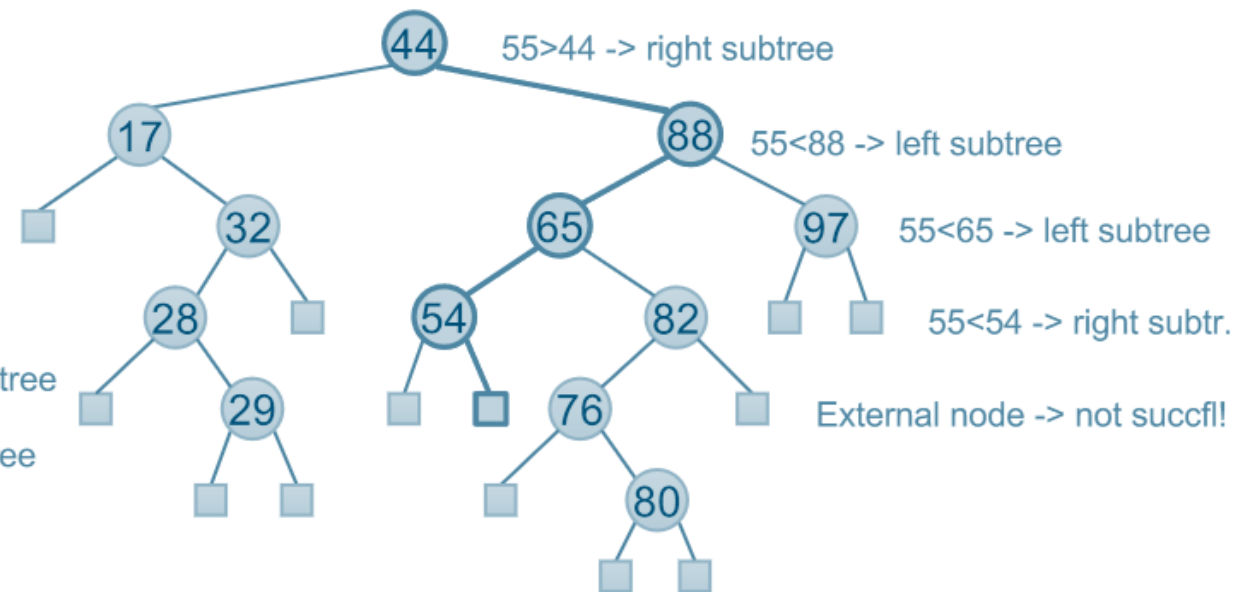
return Treesearch(k, T.rightChild(v))

Examples for Search in Binary Search Trees

TreeSearch(76, root)
(**sucessful** search)



TreeSearch(55, root)
(**unsucessful** search)



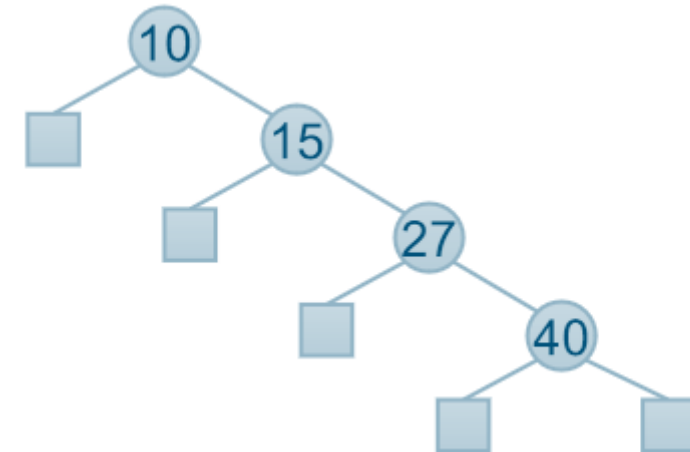
Complexity of Binary Search Trees

For **searching**, **inserting**, **removal**, nodes along a **root-leaf-path** (plus possibly siblings of such nodes) needs to be traversed.

- Complexity per node: $O(1)$
- Therefore the runtime is $O(h)$, where h is the **height of the tree**

In worst case the height of the binary search tree is $h = N$
(tree „degenerates“ to sorted sequence)

- Worst case complexity **$O(N)$**
- Therefore, the tree must be kept **balanced** to allow operations in **$O(\log N)$**



Trees

Problem: trees can degenerate (e.g. become lists)

Removal of a node

- **Very easy:** Removal of a leaf
- **Easy:** Removal of a node **p** with 1 successor:
replace node with successor
- **More difficult:** Removal of a node **p** with 2 successors:
Search for the **leftmost node q** in **right subtree** (symmetric successor).

Replace **p** with **q**, remove **q** from its original position.

Binary Search Tree :: Access Cost

Cost measure: Number of nodes visited or number of search steps or key comparisons required.

Average access cost z of a tree **B** ...

- is obtained by calculating its **total path length PL** as the **sum of the lengths of all paths** from the root to each node K_i .

$$PL(B) = \sum_{i=1}^N Level(K_i)$$

- The mean path length is calculated by $l = \frac{PL}{N}$

Maximum access cost We have the longest search path and thus the maximum access cost, when the binary search tree degenerates into a linear list.

$$height\ h = l_{max} + 1 = N$$

Binary Search Tree :: Access Cost

Maximum average access cost:

$$z_{max} = \frac{1}{N} * \sum_{i=0}^{N-1} (i + 1) * 1 = N * \frac{(N + 1)}{2N} = \frac{(N + 1)}{2} = O(N)$$

Minimum average access cost:

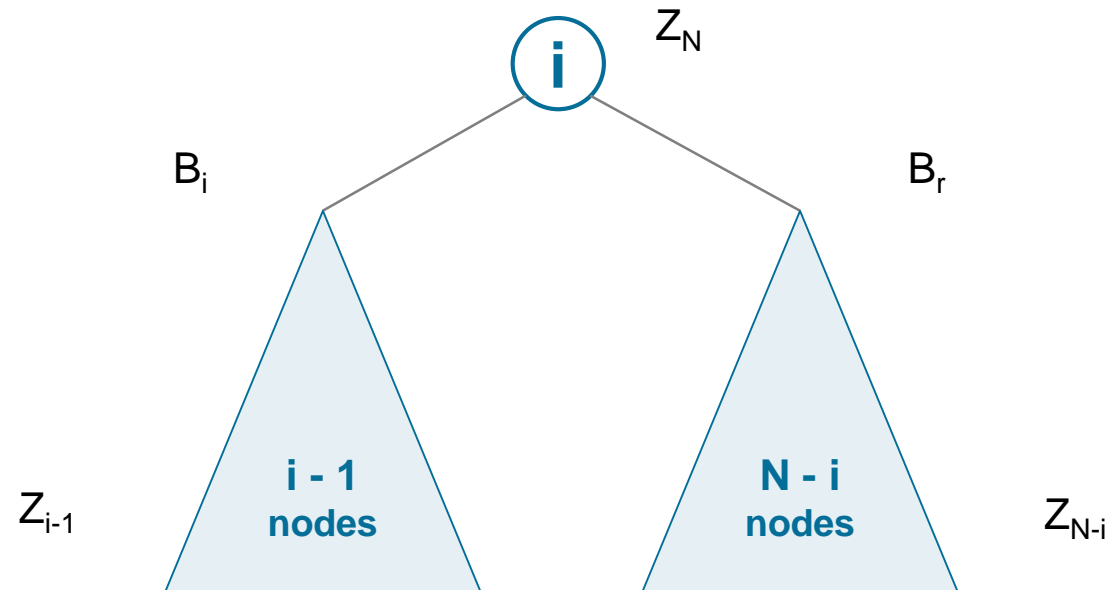
- They can be expected in an **almost complete** or **balanced** tree structure.
- Total number of nodes: $2^{h-1}-1 < N \leq 2^h-1$
- height $h = \lceil \log_2 N \rceil + 1$
- Minimum average access cost: $z_{min} \approx \log_2 N - 1$

Binary Search Tree :: Access Cost

Average access cost

- Extreme cases of average access costs have only little significance.
- The **difference** between the **average** and the **minimum** access cost is a measure of the **urgency of balancing**.

Determination of the average access cost:



Binary Search Tree :: Access Cost

N different keys with values 1, 2, ..., N are given in random order.

The probability that the first key has the value i is $\frac{1}{N}$
(Assumption: same access probability to all nodes)

For the tree with i as root we get:

$$Z_N(i) = \frac{1}{N} * ((Z_{i-1} + 1) * (i - 1) + 1 + (Z_{N-i} + 1) * (N - i))$$

The recursive equation can be represented in non-recursive, closed form by means of the [harmonic function](#):

$$H_N = \sum_{i=1}^N \frac{1}{i}$$

Binary Search Tree :: Access Cost

We get:

$$z_N = 2 * \frac{(N + 1)}{N} * H_N - 3 = 2 \ln(N) - c$$

Relative additional cost:

$$\frac{z_N}{z_{min}} = \frac{2 \ln(N) - c}{\log_2(N) - 1} \sim \frac{2 \ln(N) - c}{\log_2(N)} = 2 \ln(2) = 1,386 \dots$$

The **balanced binary search tree** causes the **least cost** for all basic operations.

However, **perfect balancing** at any time is very **expensive**.

Balanced Trees

Efficiency of dictionary operations (**insert**, **search**, **remove**) on trees depends directly on the **tree height**.

Tree with N nodes:

- Minimum height: $\lfloor \log_2 N \rfloor$, maximum height: $N-1$
- Access on average $O(\log_2 N)$, but in worst case linear.

Aim:

- Fast access with $z_{max} \sim O(\log_2 N)$
- Insert and remove operations in logarithmic complexity.

Heuristics:

For each node in the tree, the number of nodes in each of its two subtrees should be kept as constant as possible.

Balanced Trees

Two different approaches

- **Height-balanced trees:**

The maximum allowed **height difference** of the two subtrees is limited

- **Weight-balanced trees:**

The **ratio** of the **node weights (number of nodes)** of both **subtrees** meets certain conditions.

Height-Balanced Search Trees

- **B-Trees**

always **balanced** due to **balance sustaining search, insert, remove**
number of children per inner node between t ($=m/2$) and $2t$ ($=m$)

access in $O(\log(n))$

- **AVL Trees**

The heights of **each internal node's children** only **differ** by a **maximum of 1**

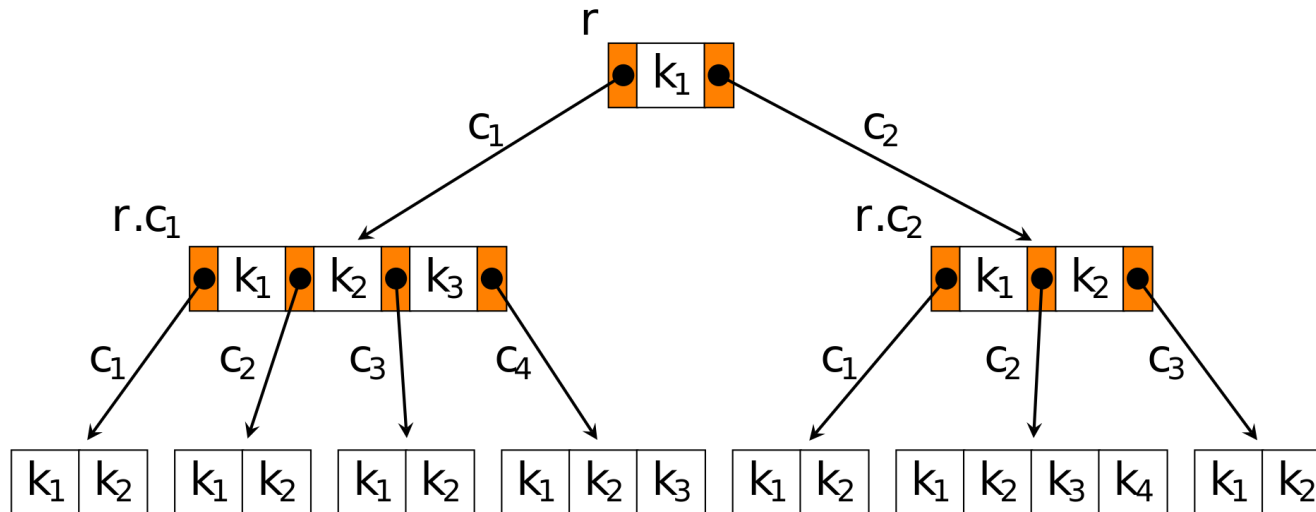
access in $O(\log(n))$

Example from <https://de.wikipedia.org/wiki/B-Baum>

B-Trees

Tree structure (m-Tree) that is **always balanced**

- **self-balancing** mechanism integrated in access operations (**search, insert, remove**)
- maintains **sorted** data (**keys**)
- **search()**, **insert()**, **remove()** in **logarithmic time**
- supports "**locality**"-principle in memory organization



Graphics: Flying sheep, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=19226668>

1. Every node has at most ***m*** children (***m*** = 2t).
2. Every non-leaf node (except root) has at least $\lceil t = \frac{m}{2} \rceil$ child nodes.
3. The root has at least 2 children if it is not a leaf node.
4. A non-leaf node with ***k*** children contains ***k*** - 1 keys.
5. All leaves appear in the same level (height *h*)

B-Trees

m = 2t ... **maximum** number of children per node
n ... number of keys stored,

note:
in algorithm examples below
(t-1) – (2t-1)

then the **height h** is: $h \leq \log_t \left(\frac{n+1}{2} \right) + 1$

(therefore access in **O(log(n))**)

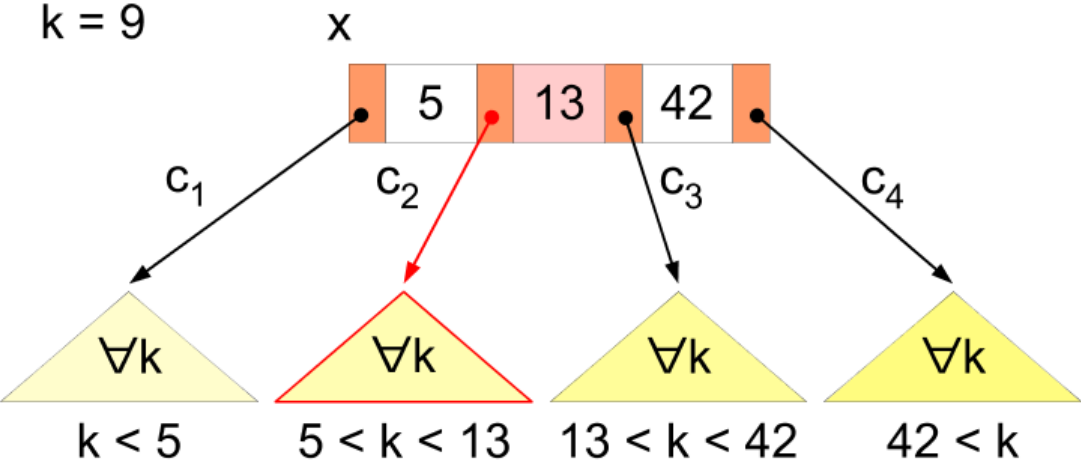
t ... **minimum** number of children per node
n ... **t^h-1**

let **t = 1024, h = 4** then

$1024^4 - 1 = (2^{10})^4 - 1 = 2^{40} - 1 = 1.099.511.627.775$ keys (i.e. access time 10 times lower)

B-Trees

Search (key = 9)



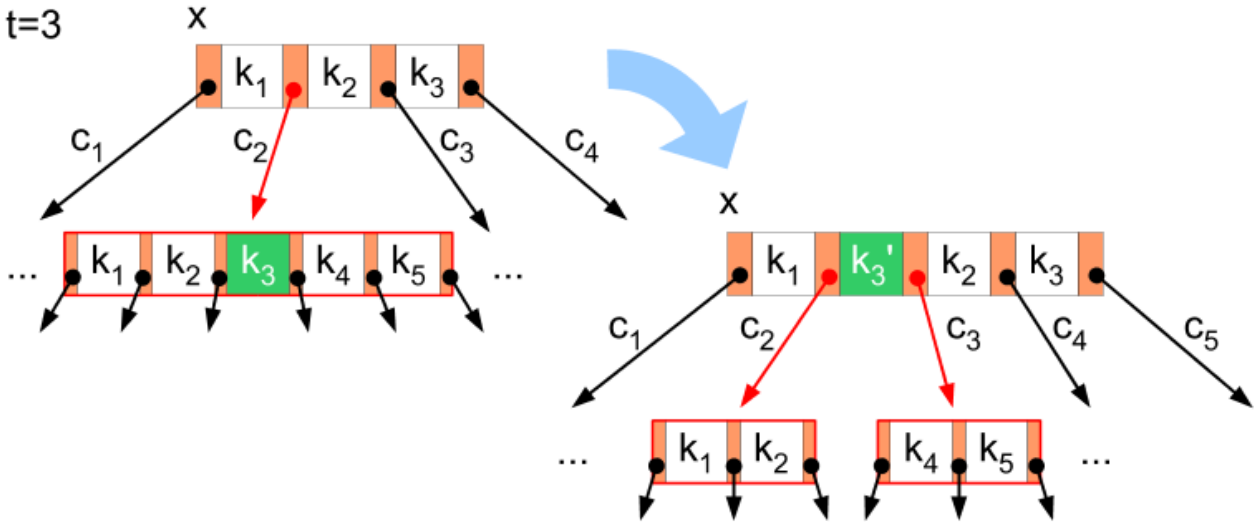
Insert

Search for insert position

(inner node)

If inner node **full**

(i.e more than $2t-1$ keys)
then preventively **split**
before stepping down



Graphics: Haui, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=245051>, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=245226>

B-Trees

Remove

Removing (**inner node**) could
destroy balance

before stepping down

check whether there are
enough keys in the **subtree**

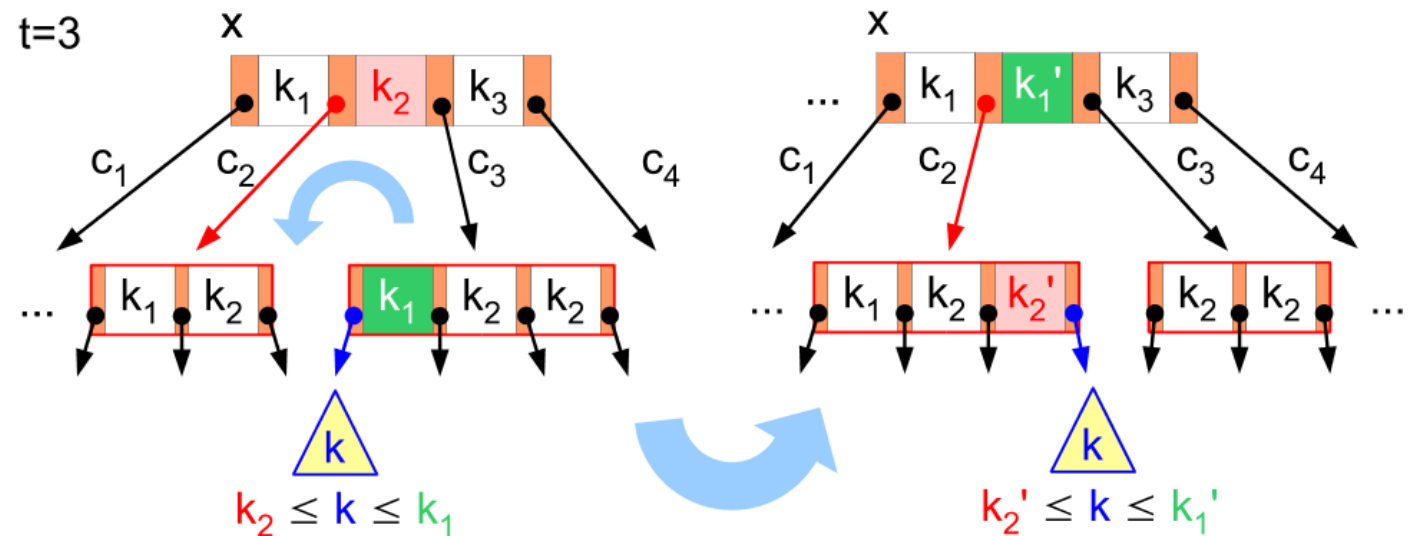
If considered inner node in
subtree has just minimum
number of keys (**t-1** keys), then

Move / Merge

to prevent from imbalance after
Remove

Move

transfer a sibling
to **expand (minimum) inner node**



B-Trees

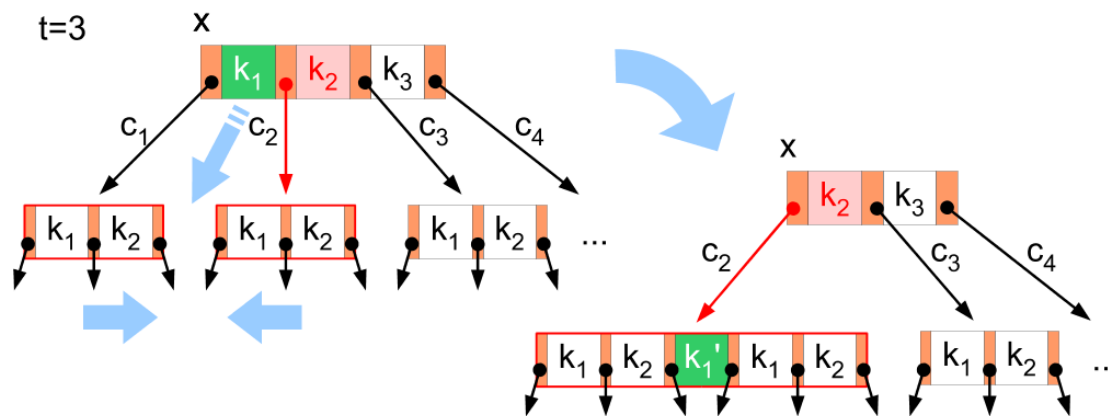
Merge

before stepping down

merge two a siblings

if both (left and right child)

are at **minimum number of inner nodes**



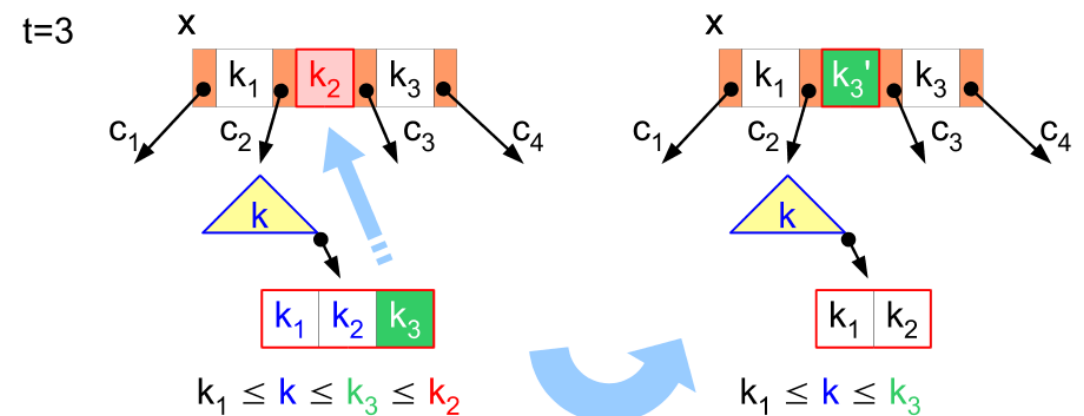
Remove

before deleting key from inner node

separate/adjust the **key ranges**

of both left and right child

(here: inorder transfer,
i.e. rightmost of left child)



Graphics: Haui, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=248208>, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=248209>

B-Trees

Example ($t=2$) i.e. min 1, max 3 keys per inner node

a–c) **insert** 5, 13, 27

d–e) **insert** 9 – induces **splitting** of root

f) **insert** 7 leaf node

g–h) **insert** 3 needs **splitting**

i–j) **remove** 9, needs induces **move** of sibling

k–l)

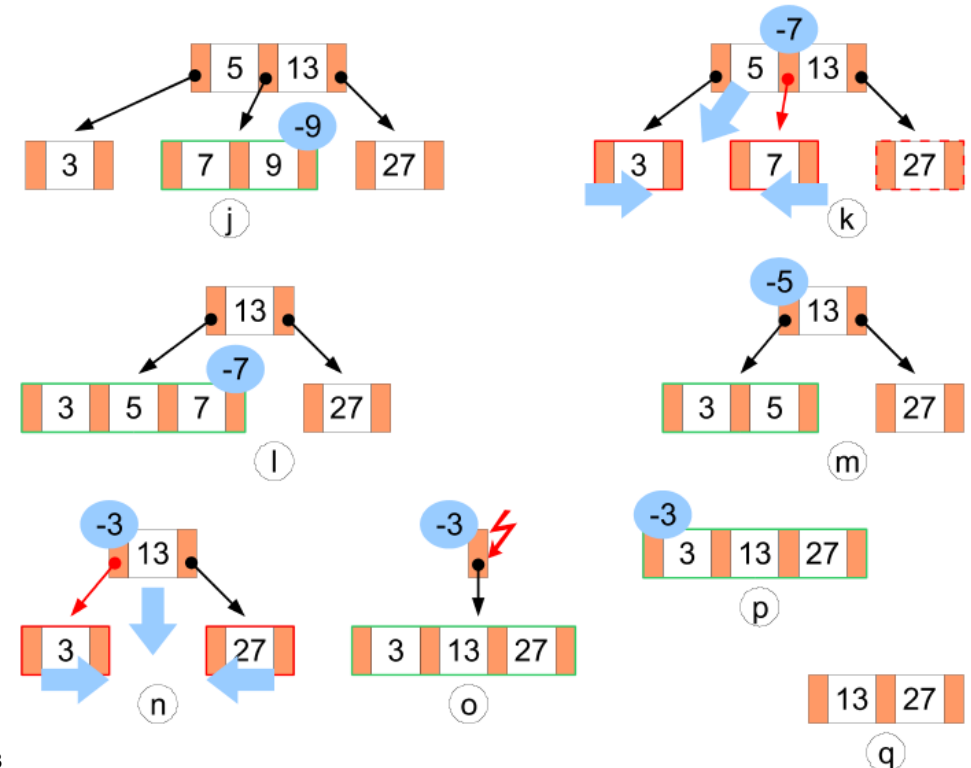
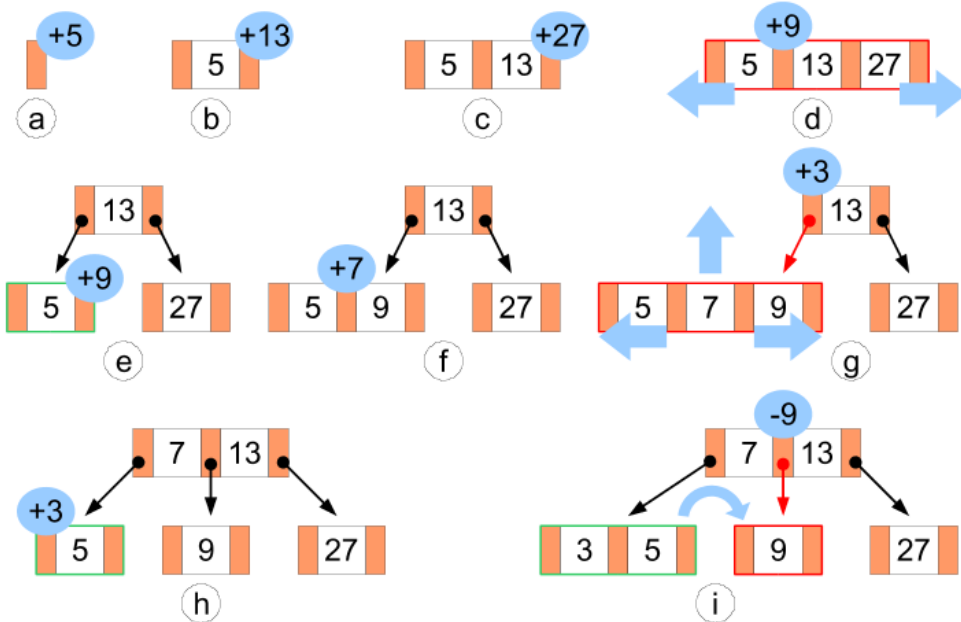
remove 7 induces **merge**

m)

remove 5 (**leaf**)

n–q)

remove 3 induces **merge** of two children of root
empty root replaced by only child



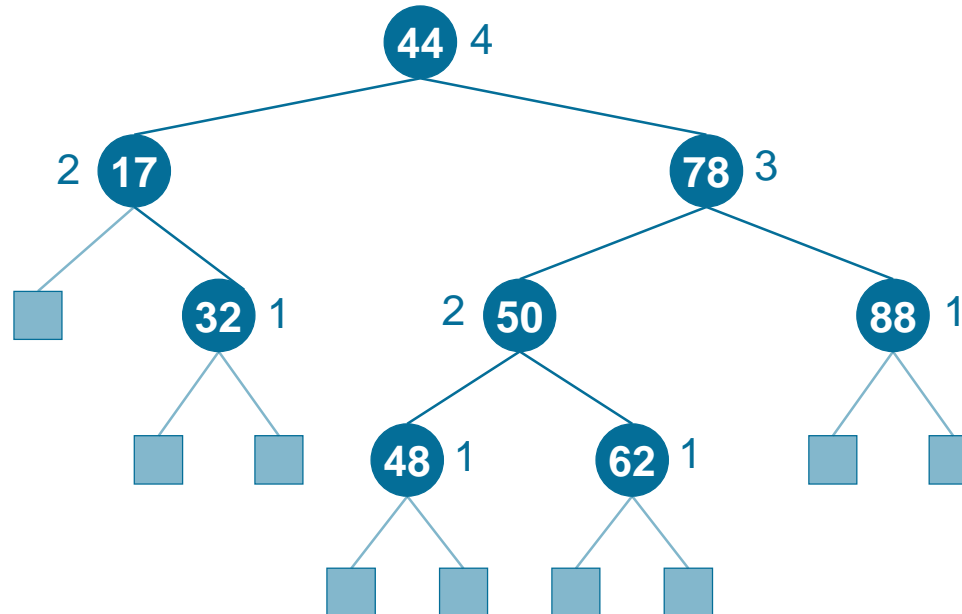
Example from <https://de.wikipedia.org/wiki/B-Baum>; Graphics: Haui, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=250188>

AVL Trees (Adelson-Velskii, Landis 1962)

AVL trees are **balanced binary search trees**

An AVL tree is a binary search tree, in which the **heights of each internal node's children only differ by a maximum of 1**.

Example (numbers next to the nodes indicate their height):



Height of an AVL tree

Assertion: The **height** of an AVL tree T that stores n keys is $O(\log n)$

Sketch of a proof:

Find $n(h)$, the smallest possible number of internal nodes in an AVL tree of height h .

- Trivial: $n(1) = 1$, $n(2) = 2$

Height of an AVL tree

Assertion: The height of an AVL tree T that stores n keys is $O(\log n)$

Sketch of a proof:

Find $n(h)$, the smallest possible number of internal nodes in an AVL tree of height h .

- Trivial: $n(1) = 1$, $n(2) = 2$
- $n \geq 3$: The AVL tree of height h with minimum $n(h)$ consists of a root node, an AVL subtree of height $h-1$ and an AVL subtree of height $h-2$.

i.e.

$n(h) = 1 + n(h-1) + n(h-2)$
as $n(h-1) > n(h-2)$ it follows:

**Fibonacci progression
(exponential!)**

- findElement: $O(\log n)$
- findAllElements $O(\log n + s)$
(s =number of findings)

$n(h) > 2 n(h-2)$
 $n(h) > 4 n(h-4)$
 $n(h) > 8 n(h-6)$

choose $i = \lceil h/2 \rceil - 1$

...

$n(h) > 2^i n(h-2i)$ then $n(h) \geq 2^{h/2-1}$

Using the logarithm we get: $h < 2 \log n(h) + 2$

Therefore we have: The height of an AVL tree is $O(\log n)$

$\log n(h) > h/2 - 1$

Insert in AVL trees

By **inserting** a node into an AVL tree, the **height** of some nodes in this tree **changes**

Insert operation may cause the AVL tree to be **unbalanced**

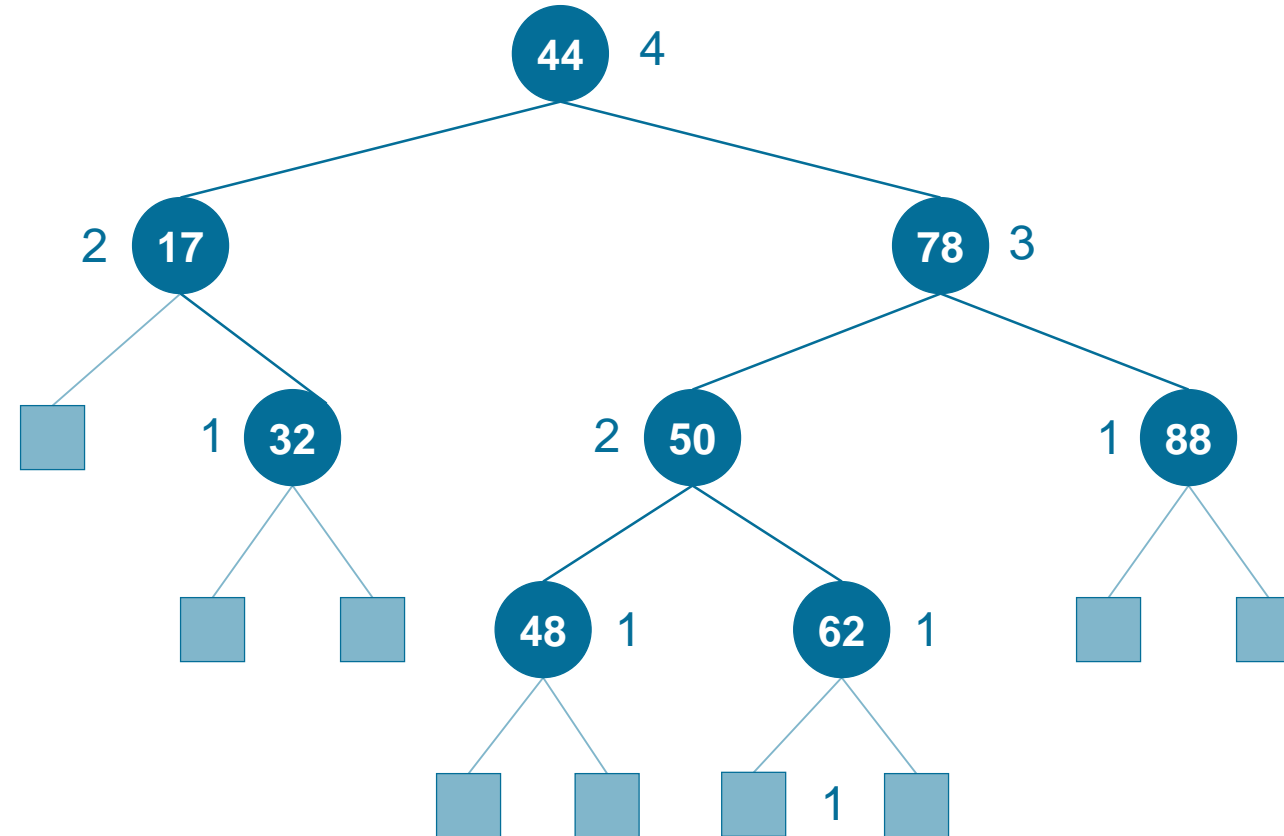
- Go up **from the new node** in the tree **until the first node x is found, whose "grandfather" z is an unbalanced node.**
- Since **z** was unbalanced by inserting a node that lies in a sub-tree with root **y** (where y is a child of **z**), we have $\text{height}(y) = \text{height}(\text{sibling}(y)) + 2$

Re-balancing of the sub-tree with root z requires **restructuring**:

- **x , y and z** are renamed to **a , b , c** (according to **in-order traversing**)
- **z** is replaced by **b** , whose children are now **a** and **c** , whose children are again **four subtrees**, which were formerly the children of x , y and z

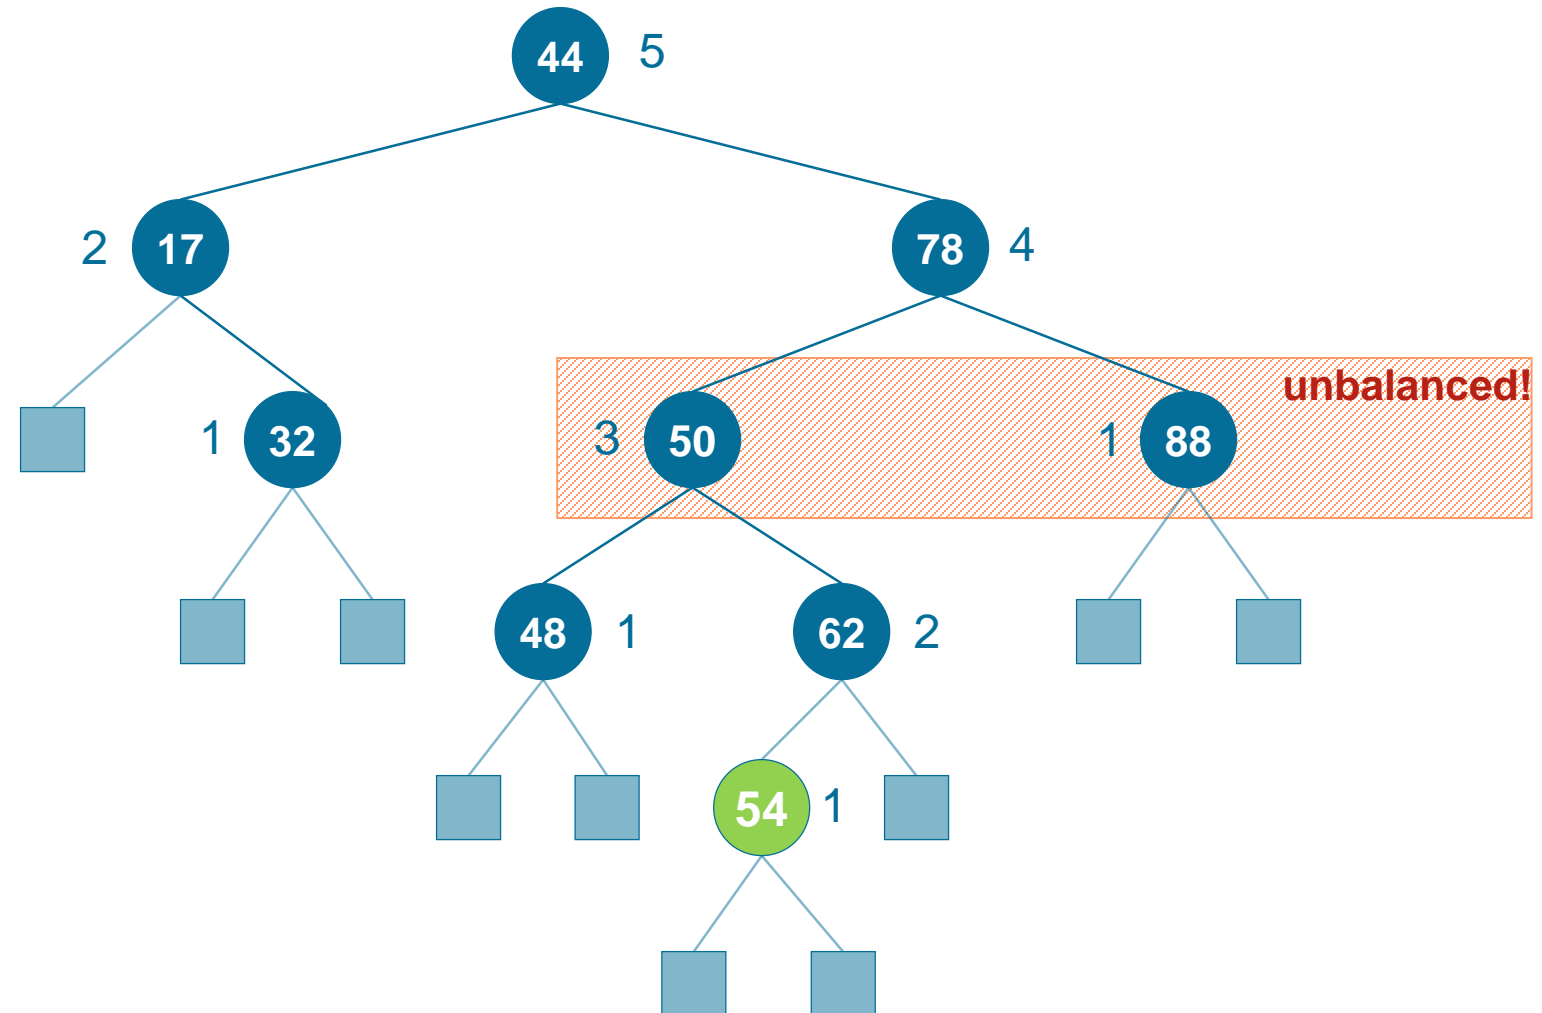
Insert in AVL trees

Insert 54



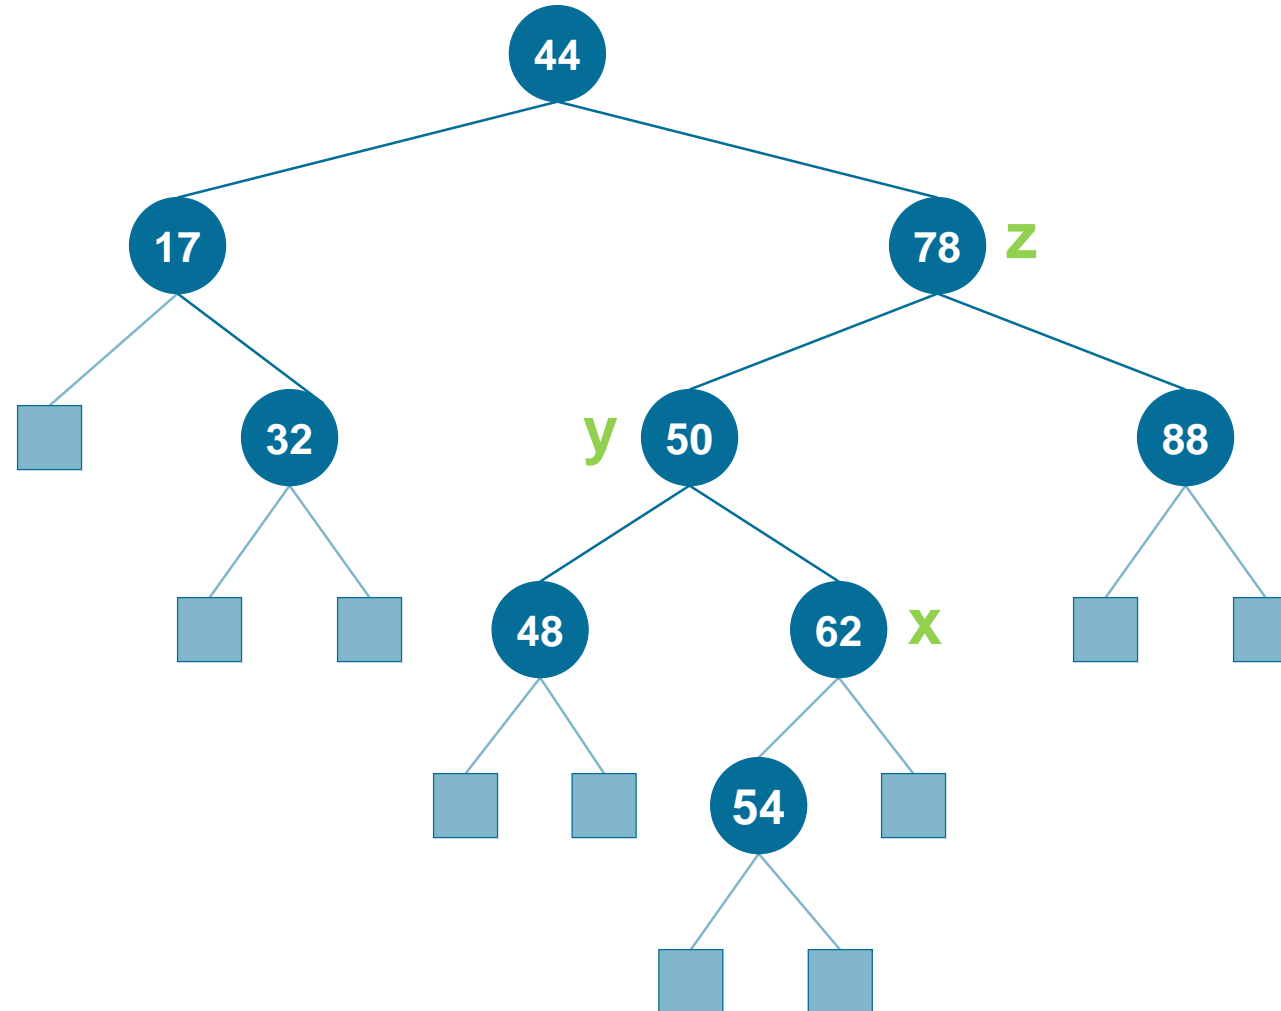
Insert in AVL trees

After inserting 54:



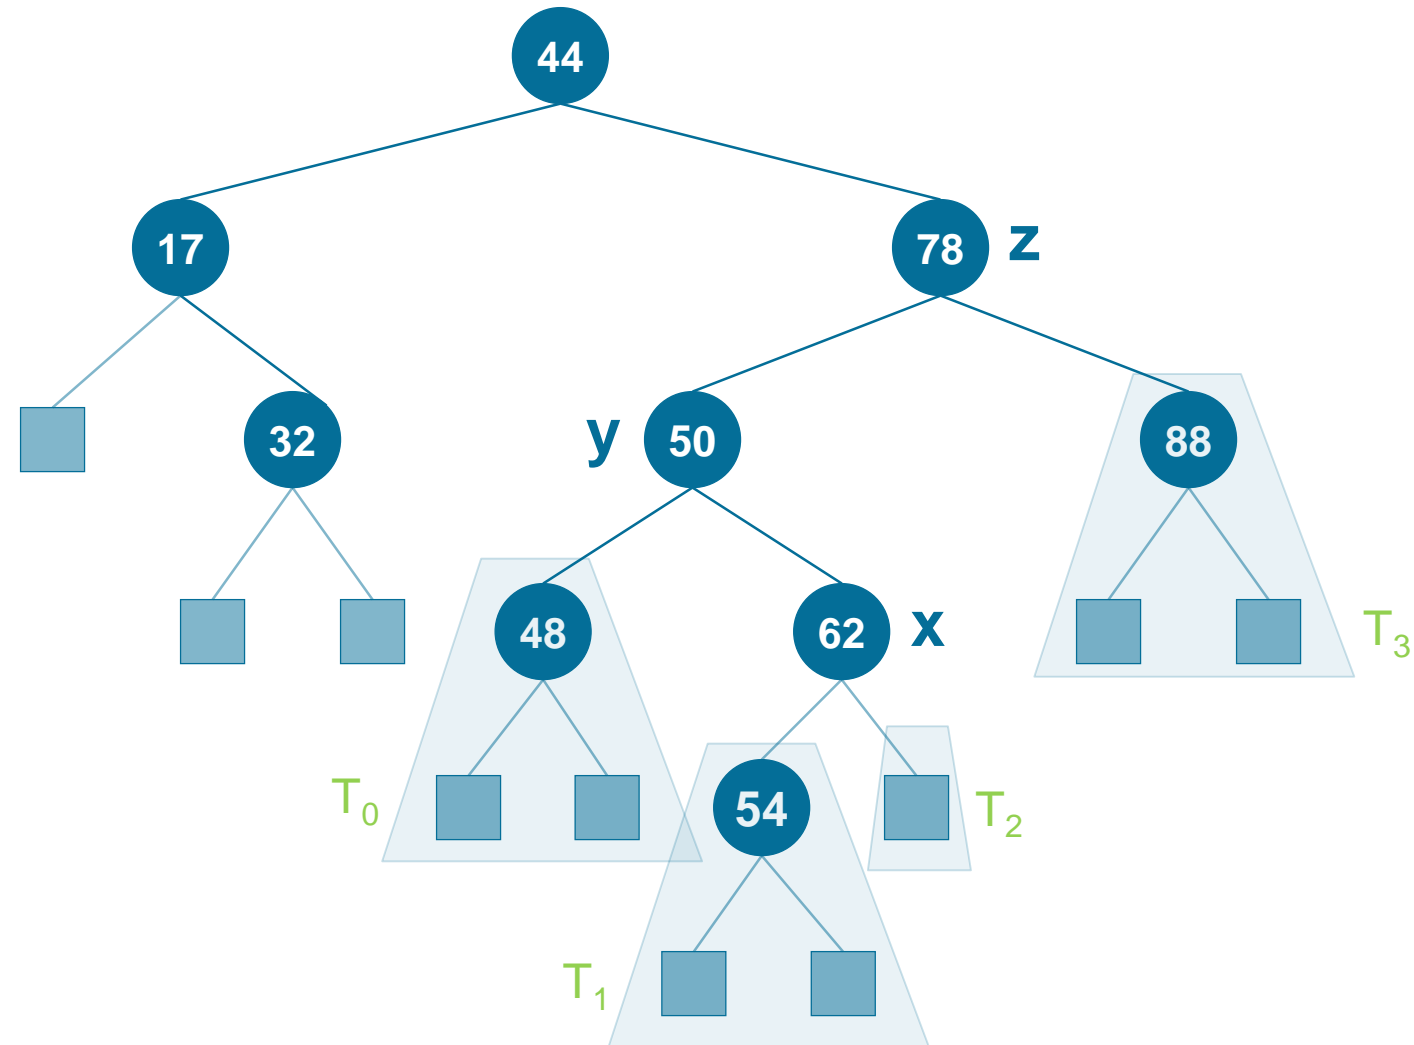
Insert in AVL trees

After inserting 54:



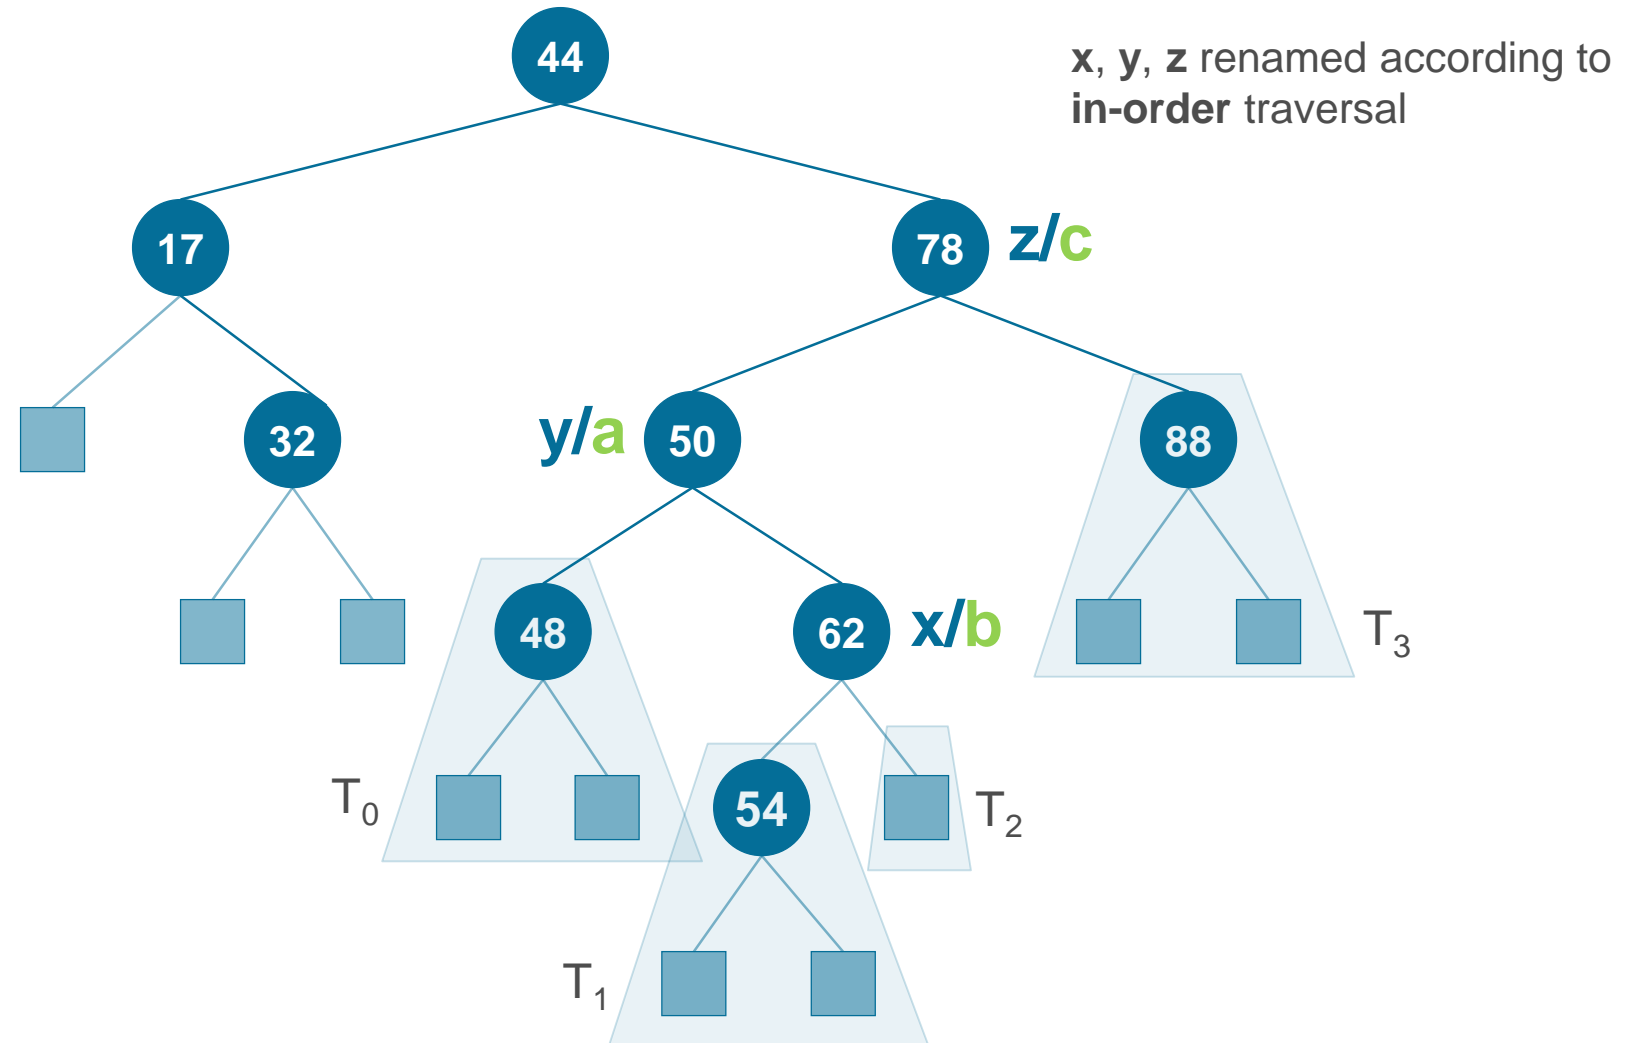
Insert in AVL trees

After inserting 54:



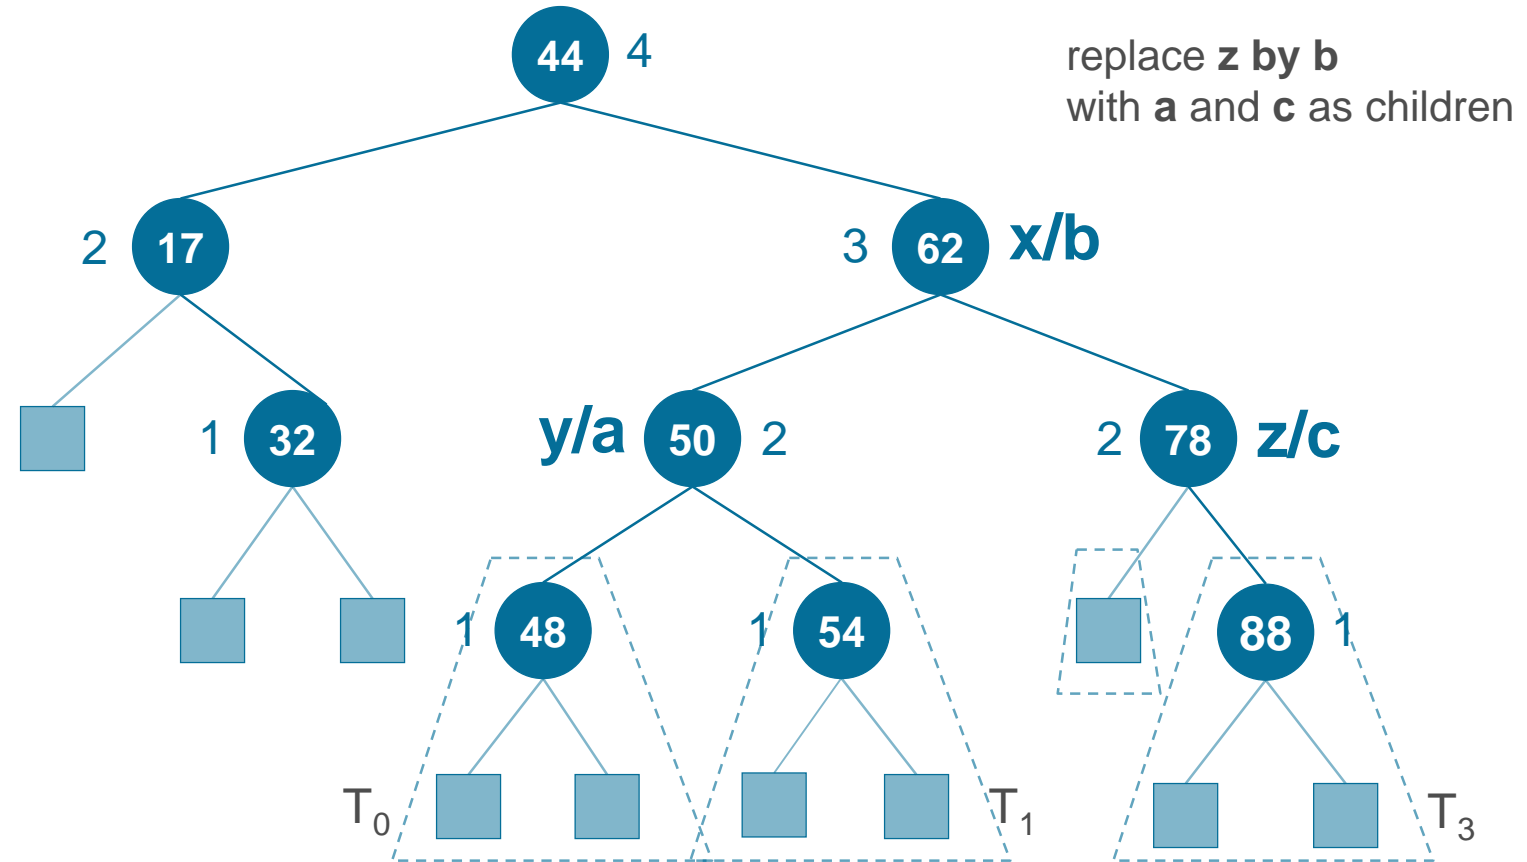
Insert in AVL trees

After inserting 54:

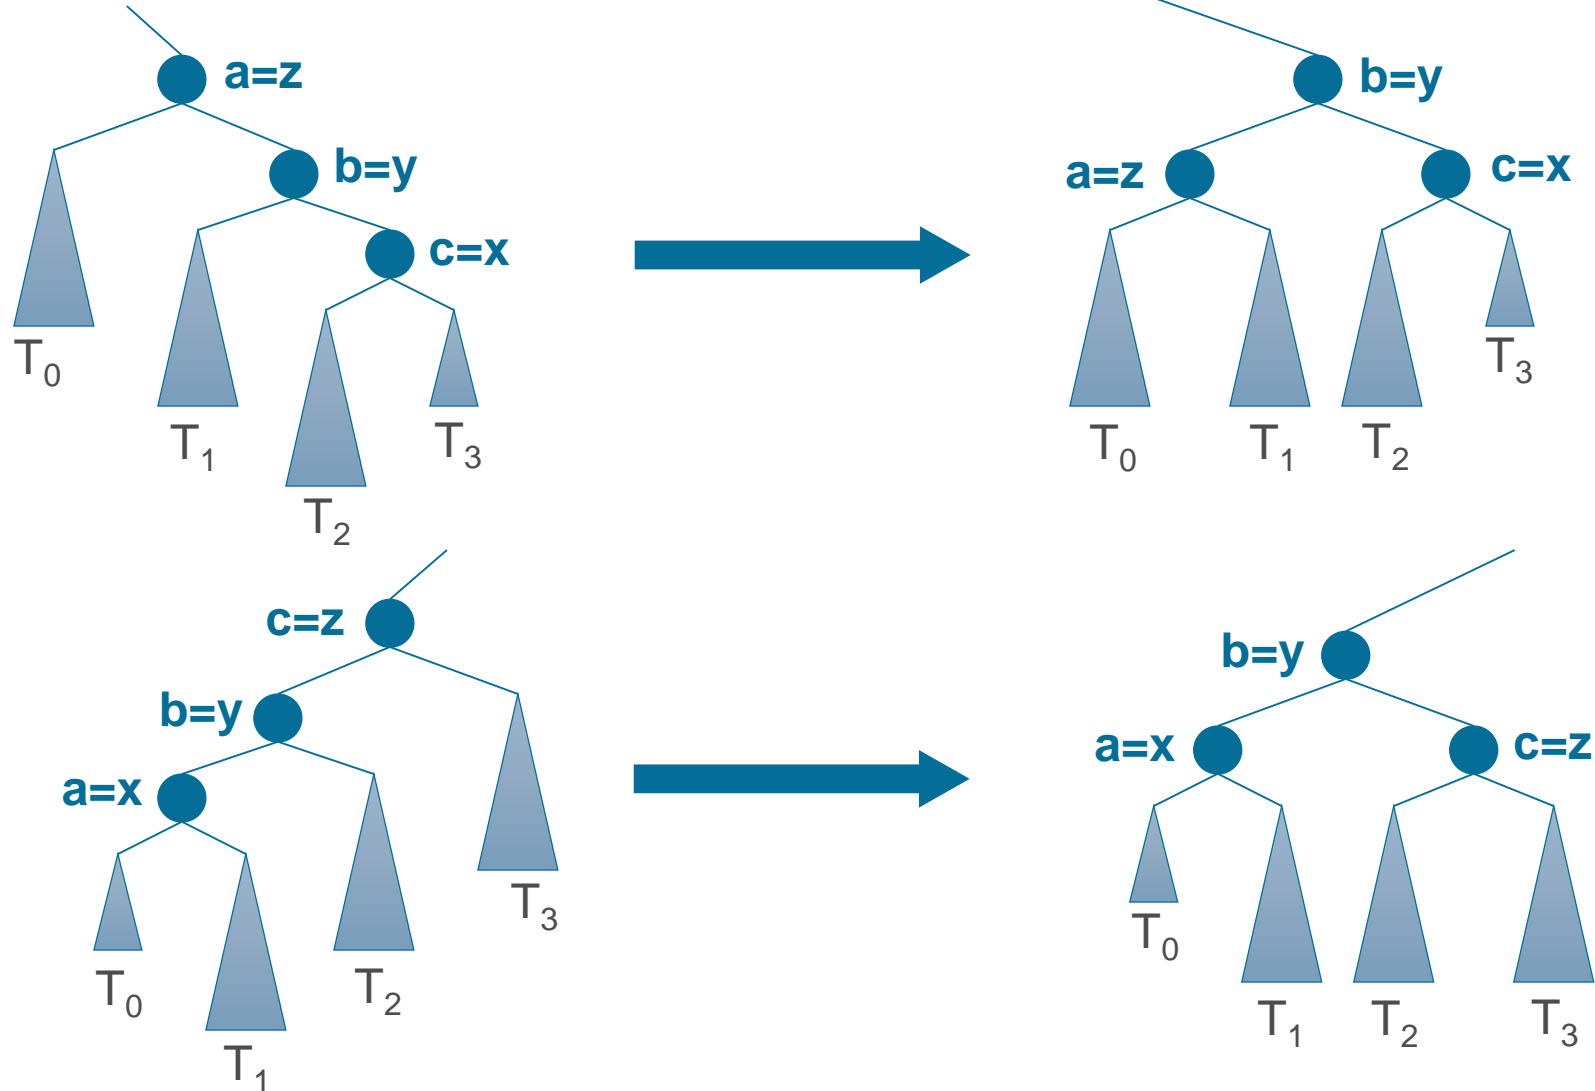


Insert in AVL trees

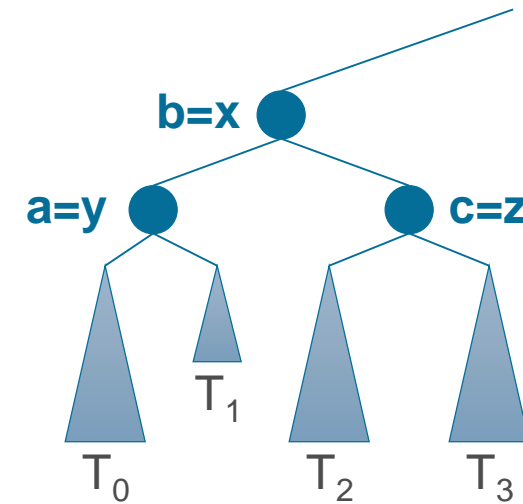
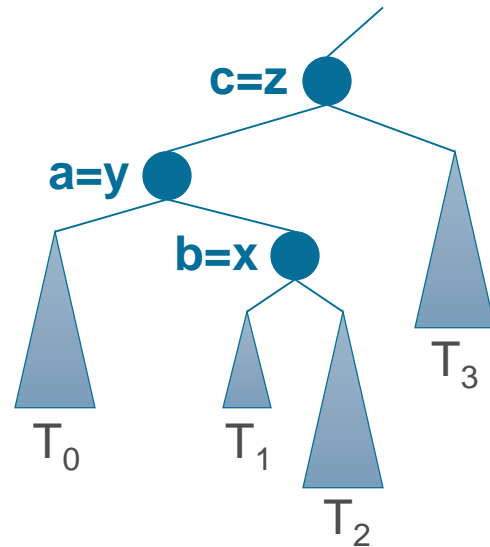
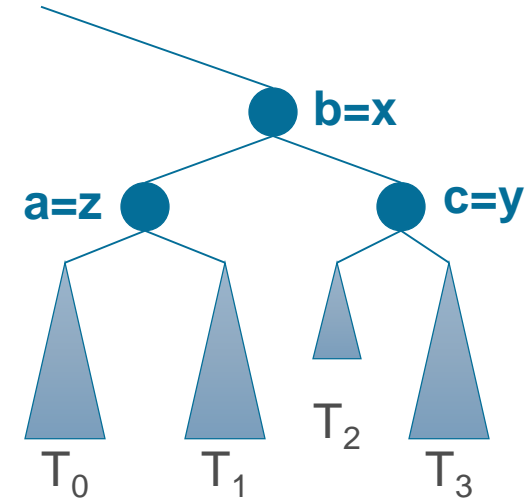
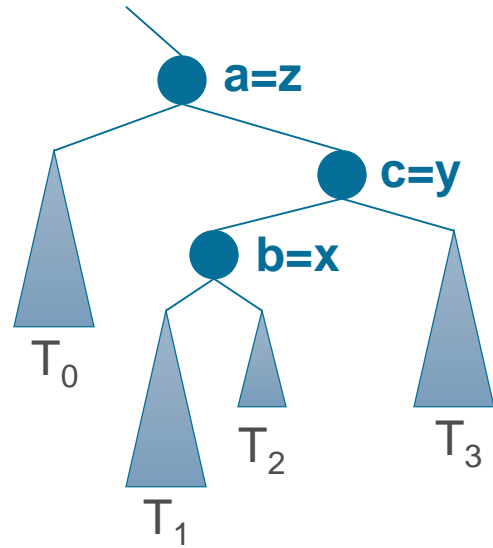
After inserting 54:



Restructuring in AVL trees (Single Rotations)



Restructuring in AVL trees (Double Rotations)



Algorithm for Restructuring

Algorithm restructure (x):

Input:

node x of a **binary tree T** with **parent y** and **grandparent z**

Output:

tree **T** **restructured** by (single or double) **rotation** of **nodes x, y, z**

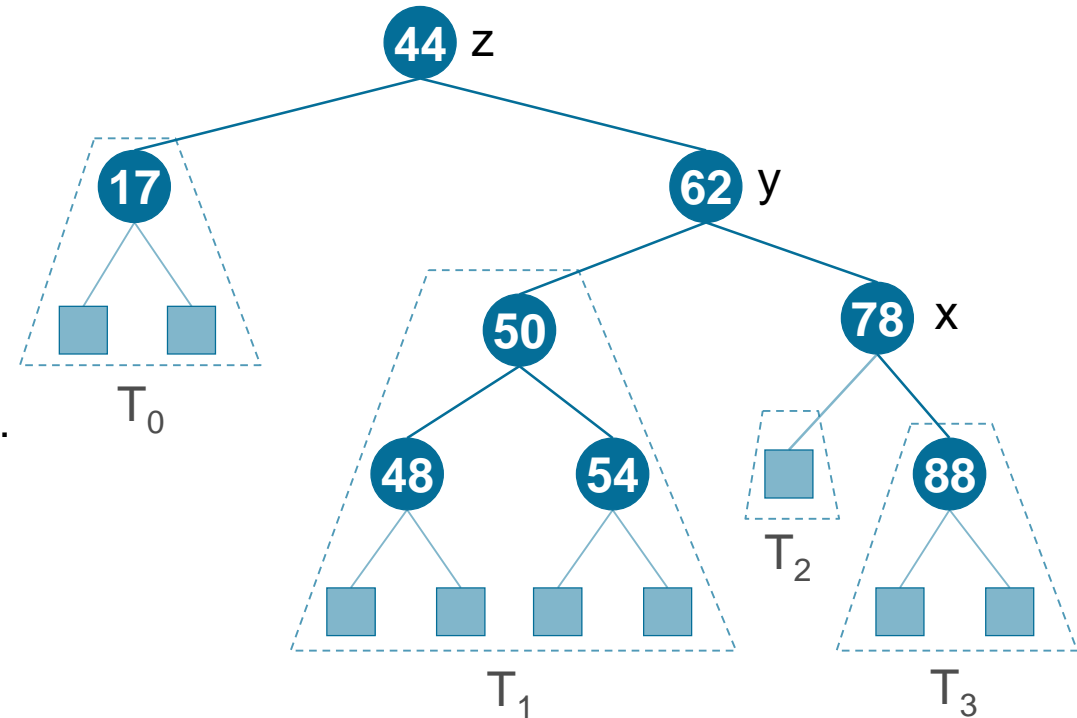
1. Let **(a,b,c)** be the **in-order-sequence** of the nodes **x, y, z** and let **(T₀, T₁, T₂, T₃)** be the in-order-sequence of the **subtrees** of **x, y** and **z**
2. **Replace** subtree with root **z** by subtree with root **b**
3. Assign **a** as **left child of b** and assign **T₀** and **T₁** as left/right **subtree** of **a**
4. Assign **c** as **right child of b** and **T₂, T₃** as left/right **subtree** of **c**

Cut/Link Restructuring Algorithm

Each tree to be balanced can be divided into **7 parts**:

- nodes **x, y, z** and
- **4 subtrees (T_0, T_1, T_2, T_3)**,
with the children of **x, y, z** as root

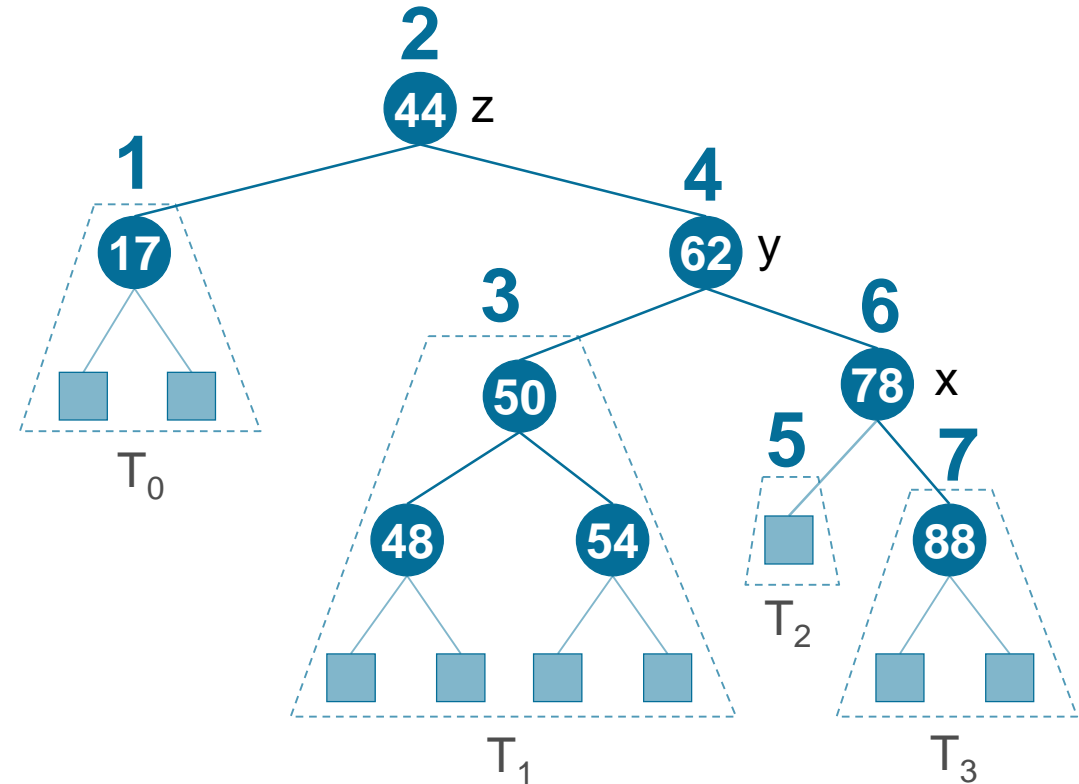
Create a **new tree** from the 7 parts, which is balanced and in which the **in-order sequence** of the parts is retained.



Cut/Link Restructuring Algorithm

Example

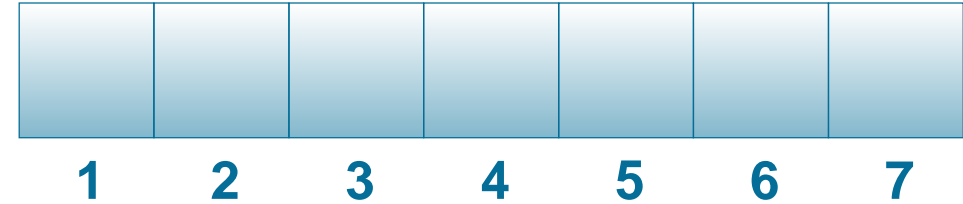
number 7 parts according to in-order traversal



Cut/Link Restructuring Algorithm

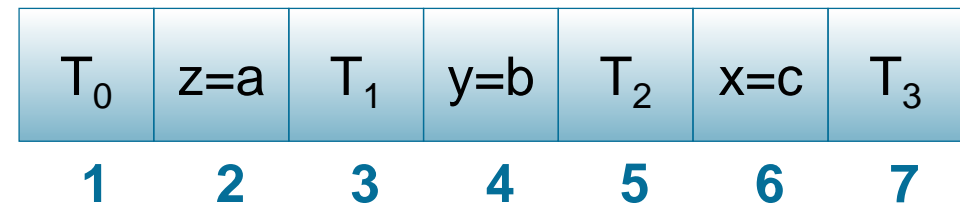
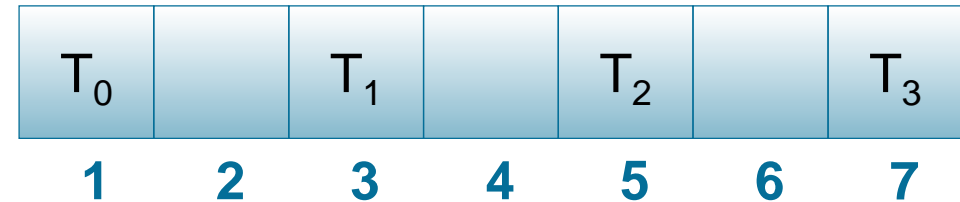
Example

Create an array with indices 1..7



“Cut” the 4 subtrees and place them into the array according to their numbering.

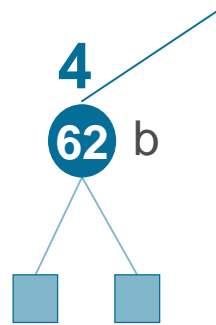
“Cut” x , y and z (in the order child, parent, grandparent) and put them in the array according to their numbering



Cut/Link Restructuring Algorithm

Example

- Reassemble the tree again
- Set element at position 4 (b) as root

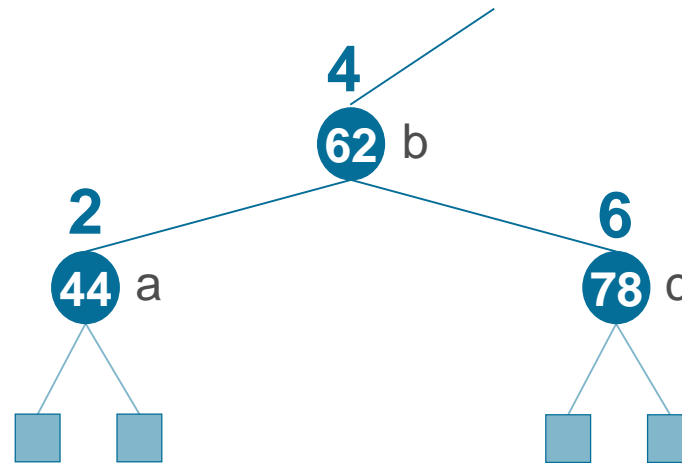


T_0	$z=a$	T_1	$y=b$	T_2	$x=c$	T_3
1	2	3	4	5	6	7

Cut/Link Restructuring Algorithm

Example

- Reassemble the tree again
- Set elements at position 2 and 6 as children

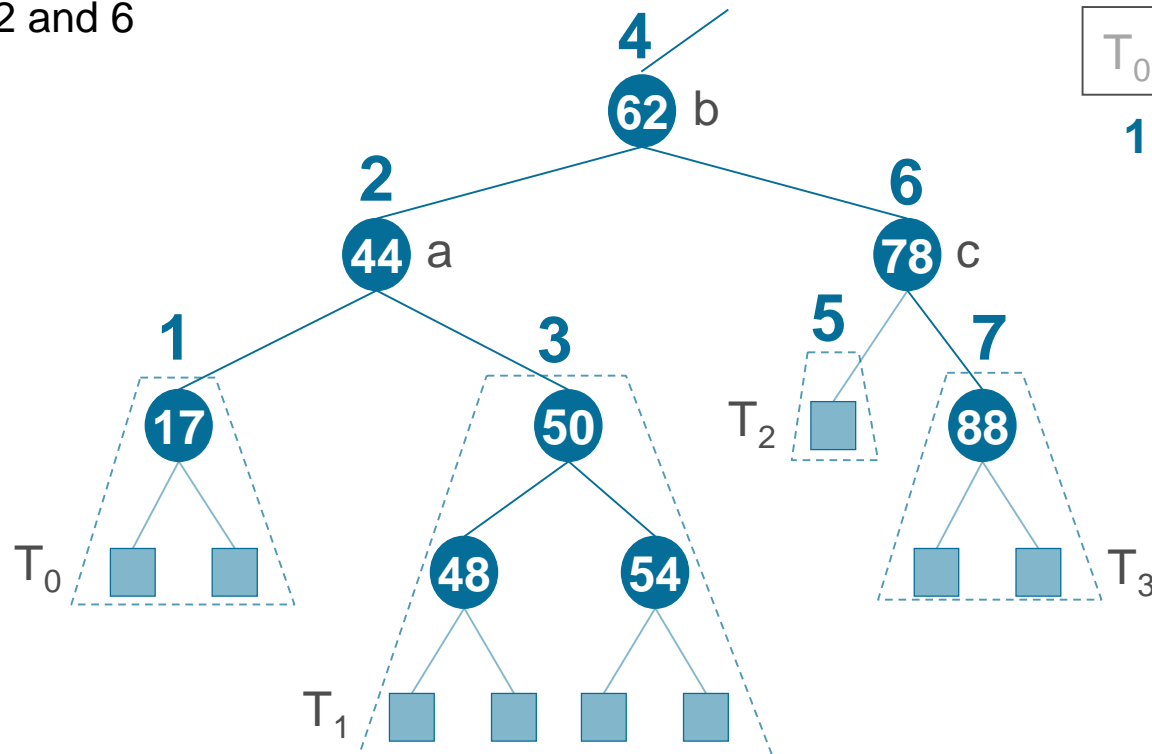


T_0	$z=a$	T_1	$y=b$	T_2	$x=c$	T_3
1	2	3	4	5	6	7

Cut/Link Restructuring Algorithm

Example

- Reassemble the tree again
- Set elements at position
1, 3 or 5, 7 as child of 2 and 6



T ₀	z=a	T ₁	y=b	T ₂	x=c	T ₃
1	2	3	4	5	6	7

Cut/Link Restructuring Algorithm

Cut/Link restructuring algorithm has the **same effects** as the **four rotation cases** previously considered

Advantage:

- **No case distinction** necessary
- More „**elegant**“ solution

Disadvantage:

- May require more code

The two procedures **do not differ** in terms of **complexity (runtime)**.

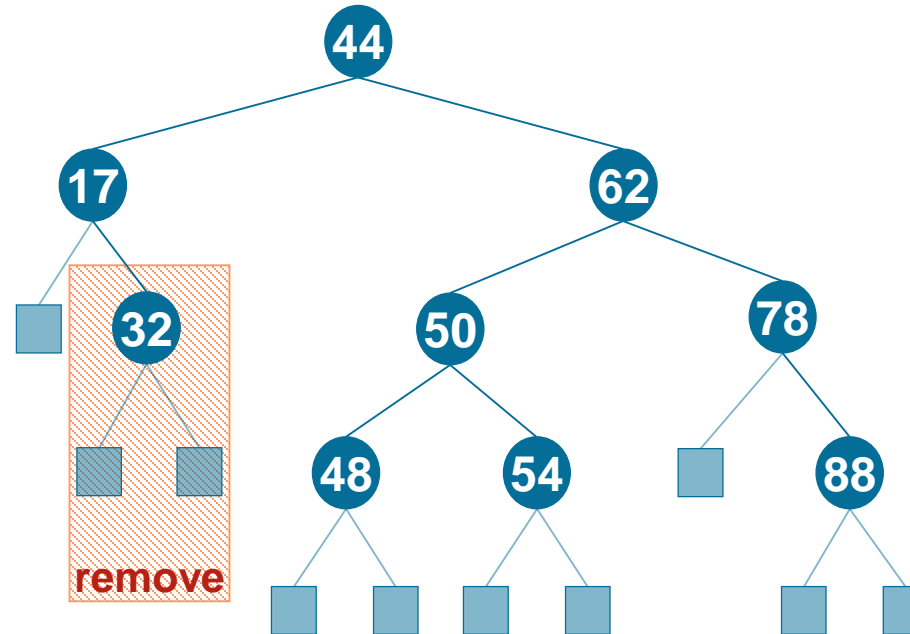
Removal of Nodes

Using `removeAboveExternal(w)` can unbalance a tree:

- Let **z** be the **first unbalanced node** which is visited while traversing **up** in the tree.
- Let **y** be the **child of z** with the **largest height** and let **x** be the **child of y** with **largest height**.
- The algorithm **restructure(x)** can be used to restructure and balance the subtree with root **z**.
- However, restructuring can **destroy the balance at higher levels**, so that the verification (and restructuring if necessary) must be **continued** until the root node is reached.

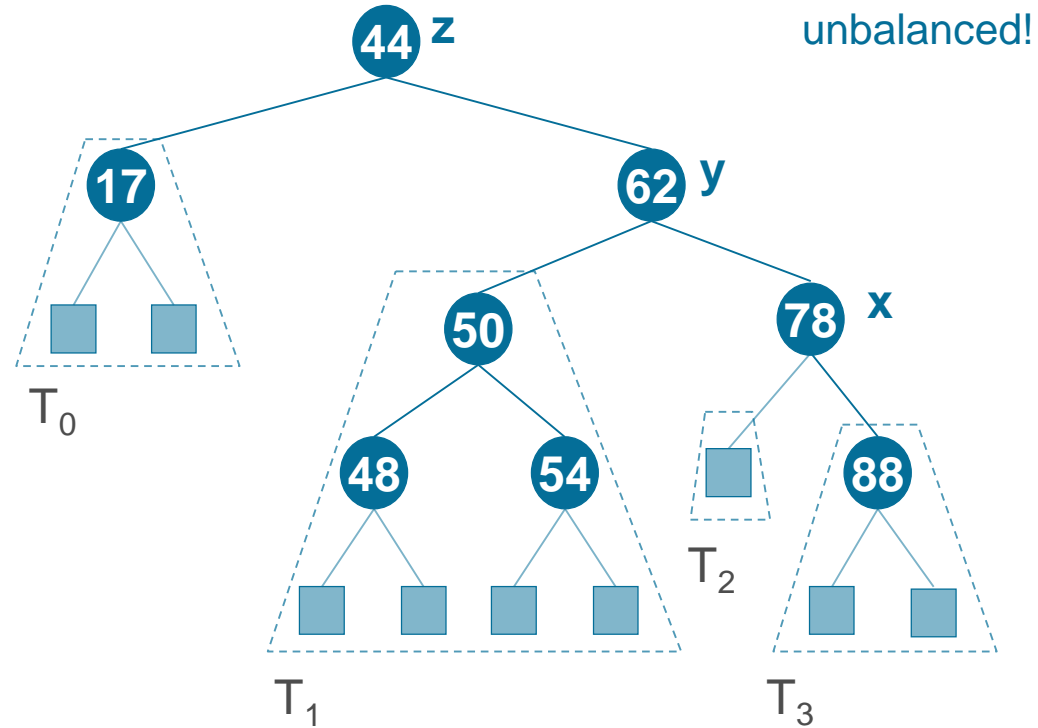
Removal of Nodes

Example



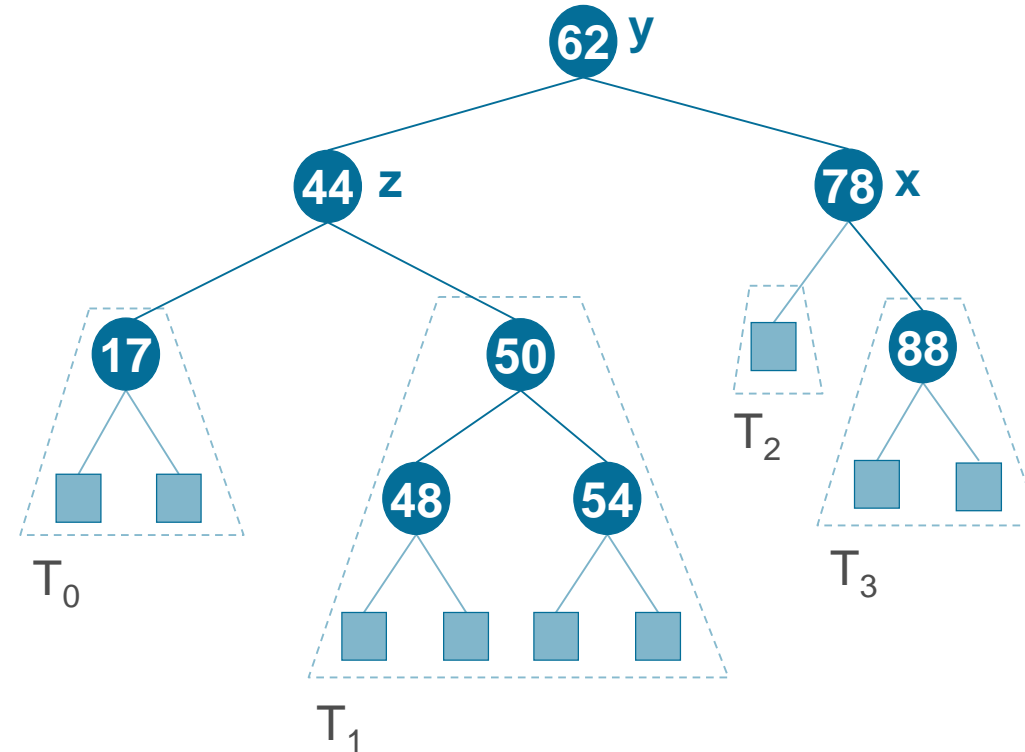
Removal of Nodes

Example



Removal of Nodes

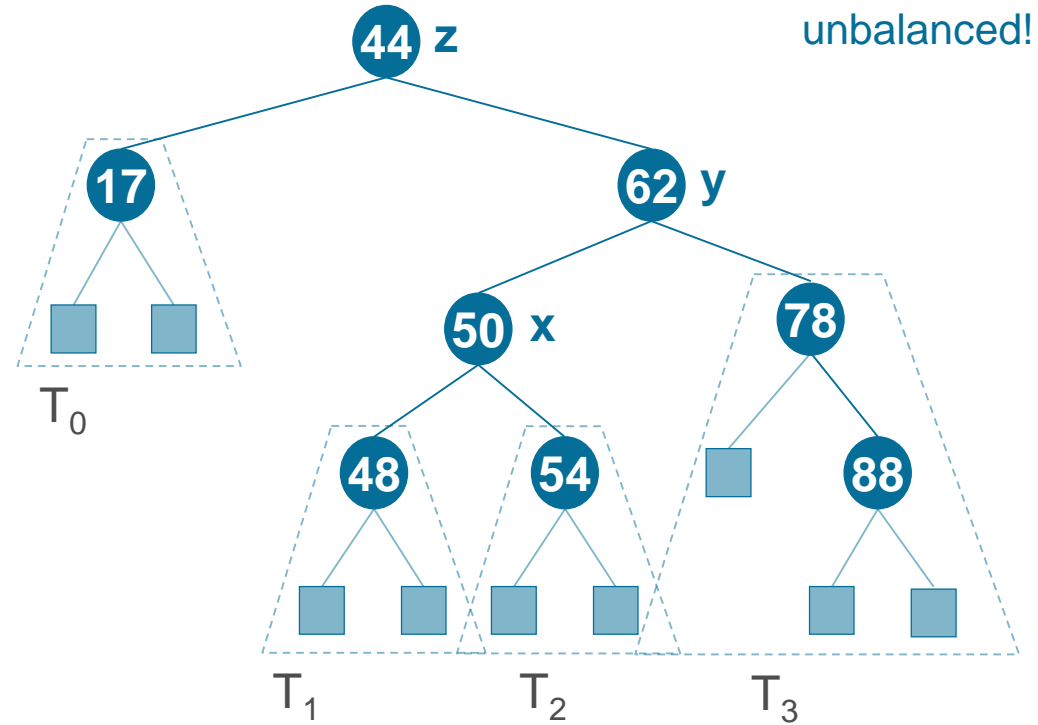
Example



Removal of Nodes

Example

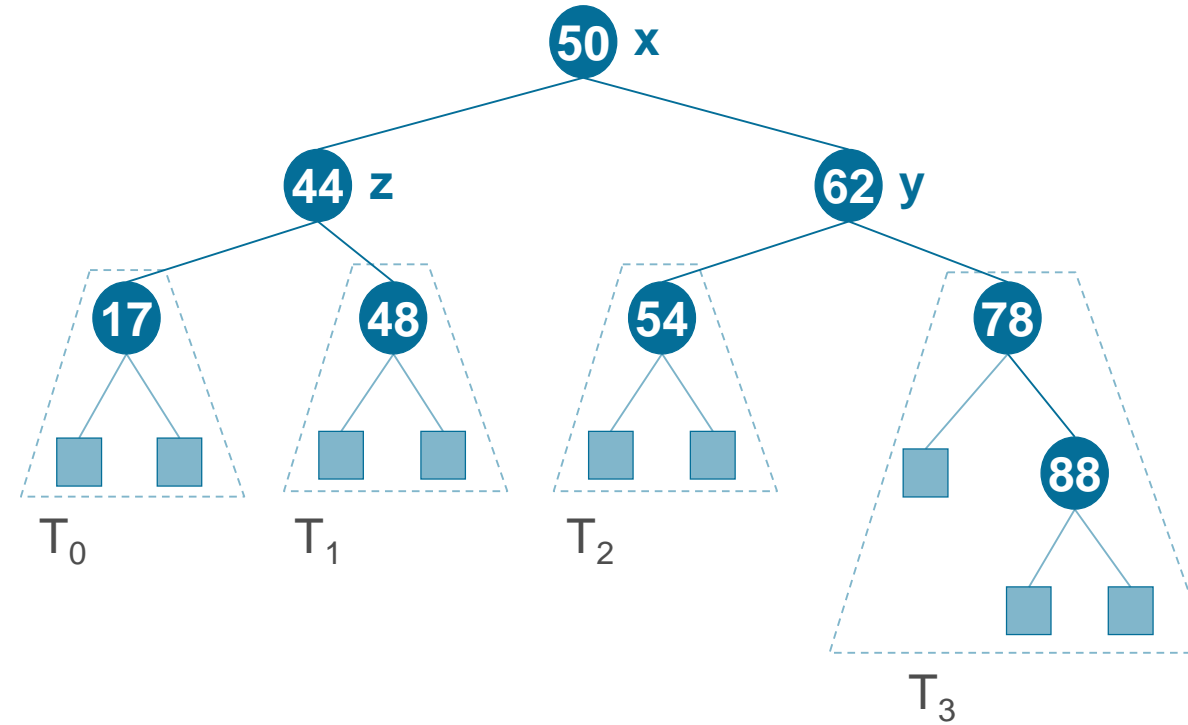
Alternative



Removal of Nodes

Example

Alternative



AVL Trees :: Complexity

Possible complexity: N = number of keys

<code>size()</code> , <code>isEmpty()</code>	$\Theta(1)$
<code>get()</code> , <code>insert()</code> , <code>remove()</code>	$O(\log N)$
<code>getAll()</code> , <code>removeAll()</code>	$O(\log N + s)$

(s = Number of elements in respective `enumeration`)

Balanced Trees

Balanced trees are introduced as a compromise between balanced and natural search trees, whereby **logarithmic search complexity** is required in the **worst case**.

For the height h_b of an AVL tree with N nodes we have:

$$\lfloor \log_2 N \rfloor \leq h_b \leq 1,44 * \log_2 (N+2)$$

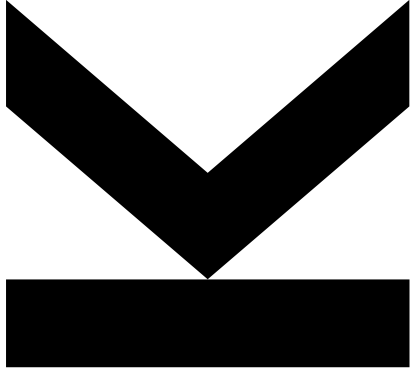
- The upper limit can be derived from Fibonacci trees, a subclass of the AVL trees.
- Let $N(h)$ be the minimum number of nodes of a height-balanced tree with height h . We have:
 - $N(0)=1, N(1)=2, N(2)=4, N(3)=7, N(4)=12, N(5)=20, \dots$
 - $N(h) = 1 + N(h-1) + N(h-2) = \text{Fib}(h+3) - 1$
 - $\text{Fib}(h) = 1/\sqrt{5} * ((1 + \sqrt{5})/2)^h - ((1 - \sqrt{5})/2)^h$
 - for all h we have: $\text{Fib}(h) \geq 1/\sqrt{5} * ((1 + \sqrt{5})/2)^h - 1$
 - if $N(h)=\text{Fib}(h+3)-1$ we have: $\log_2 (N(h)+2) \geq \log_2 (1/\sqrt{5}) + (h+3) \log_2 ((1 + \sqrt{5})/2)$
- From this the estimation follows: $h \leq 1,44 \log_2 (N(h)+2) \rightarrow h = O(\log N(h))$

0,1,1,2,3,5,8,13,21, ...

Minimum number of nodes grows exponential with height

→ so vice versa: height grows logarithmically with node number

Trees (Height Balanced)



Algorithms and Data Structures 2, 340300
Lecture – 2023W
Univ.-Prof. Dr. Alois Ferscha, teaching@pervasive.jku.at